

PMPP 2015/16



TECHNISCHE
UNIVERSITÄT
DARMSTADT

CUDA Programming (1)



(Preliminary) Course Schedule



TECHNISCHE
UNIVERSITÄT
DARMSTADT

you are here



12.10.2015	Introduction to PMPP
13.10.2015	Lecture CUDA Programming 1
19.10.2015	Lecture CUDA Programming 2
20.10.2015	Lecture CUDA Programming 3
26.10.2015	Introduction Final Projects, Exercise 1 assigned
27.10.2015	Questions and Answers (Q&A)
02.11.2015	Lecture, Final Projects assigned, Ex. 1 due, Ex. 2 assigned
03.11.2015	Questions and Answers (Q&A)
09.11.2015	Lecture, Exercise 2 due
10.11.2015	Lecture
16.11.2015	Questions and Answers (Q&A)
17.11.2015	Questions and Answers (Q&A)
23.11.2015	1 st Status Presentation Final Projects
24.11.2015	1 st Status Presentation Final Projects (continued)



(Preliminary) Course Schedule

30.11.2015

01.12.2015

07.12.2014

08.12.2014

14.12.2014

15.12.2014

Christmas Break

11.01.2015 2nd Presentation Status Final Projects (1)

12.01.2015 2nd Presentation Status Final Projects (2)

...

01.02.2015

02.02.2015

08.02.2015 final Presentation Status Final Projects (1)

09.02.2015 final Presentation Status Final Projects (2)



Review: Why PMPP?





Figure 2. Processor performance improvement between 1978 and 2006 using integer SPEC [SPEC 2006] programs. RISCs helped inspire performance to improve by 52% per year between 1986 and 2002, which was much faster than the VAX minicomputer improved between 1978 and 1986. Since 2002, performance has improved less than 20% per year. By 2006, processors will be a factor of three slower than if progress had continued at 52% per year. This figure is Figure 1.1 in [Hennessy and Patterson 2007].

from [Asanovic et al. 2006]

Comparing GPU and CPU

- floating point operations per second

Theoretical GFLOP/s



Today's Topics

- CUDA “Hello World”
- CUDA programming model
 - basic concepts and data types
- CUDA application programming interface - basic
- simple matrix multiplication example
 - illustrate basic concepts and functionalities
 - performance features will be covered later

What is GPGPU ?




- General Purpose computation using GPU in applications other than 3D graphics
 - GPU (Graphics Processing Unit) accelerates critical path of application
- data parallel algorithms leverage GPU attributes
 - large data arrays, streaming throughput
 - fine-grain SIMD parallelism
 - low-latency floating point (FP) computation
- applications – see www.GPGPU.org
 - game effects (FX) physics, image processing
 - physical modeling, computational engineering, matrix algebra, convolution, correlation, sorting, ...



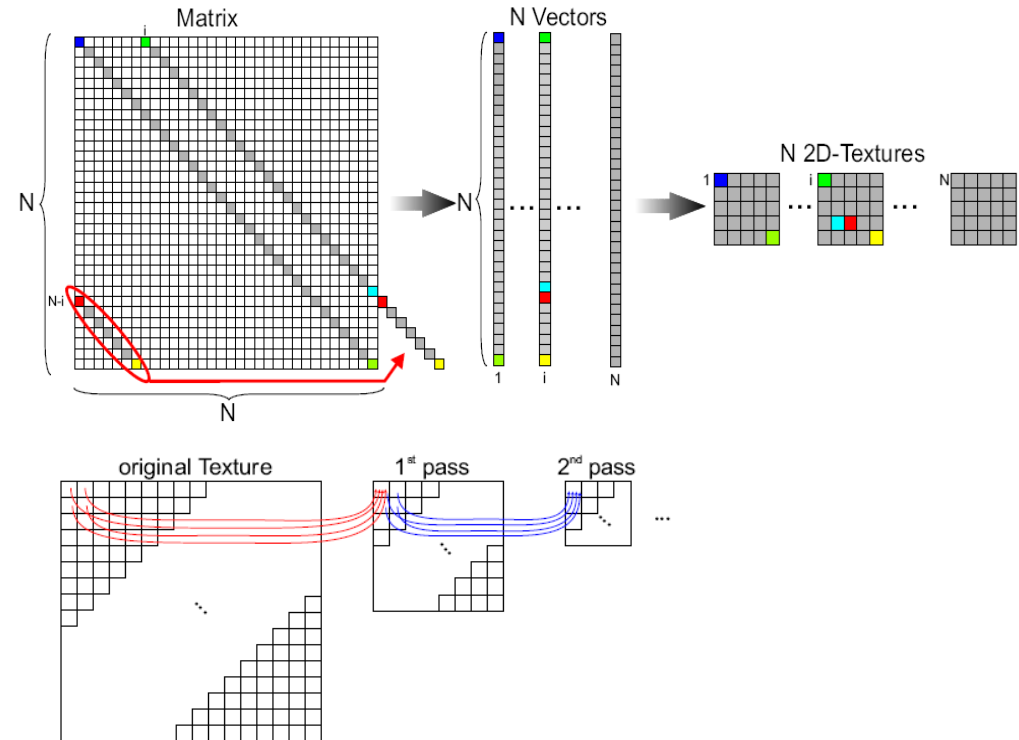
What is GPGPU ?

- traditional graphics pipeline
 - implemented as fixed-function hardware



GPGPU Example

- Jens Krüger, Rüdiger Westermann: Linear algebra operators for GPU implementation of numerical algorithms. SIGGRAPH 2003, pages 908-916, 2003.
- linear algebra on GPU
- vectors, matrices mapped to textures
- algorithms implemented as rendering passes
 - e.g., reduce operation



Typical Constraints for traditional GPGPU

- dealing with graphics API
 - working with the corner cases of the graphics API
- addressing modes
 - limited texture size/dimension
- shader capabilities
 - limited outputs
- instruction sets
 - lack of integer & bit ops
- communication limited
 - between pixels
 - scatter $a[i] = p$

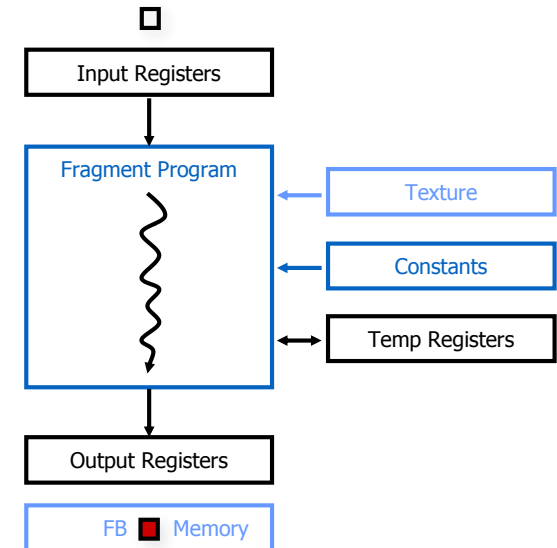




image source: NVIDIA

Modern GPU's: Unified Architecture

Discrete Design



Unified Design



Vertex shaders, pixel shaders, etc. become *threads* running different programs on flexible cores

© 2008 NVIDIA Corporation.

image source: NVIDIA

- “Compute Unified Device Architecture”
- general purpose programming model
 - user kicks off batches of threads on the GPU
 - GPU = dedicated super-threaded, massively data parallel co-processor
- targeted software stack
 - compute-oriented drivers, language, and tools
- driver for loading computation programs into GPU
 - standalone driver – optimized for computation
 - interface designed for compute – graphics free API
 - data sharing with OpenGL buffer objects
 - guaranteed maximum download & readback speeds
 - explicit GPU memory management



image source: NVIDIA



High Level View

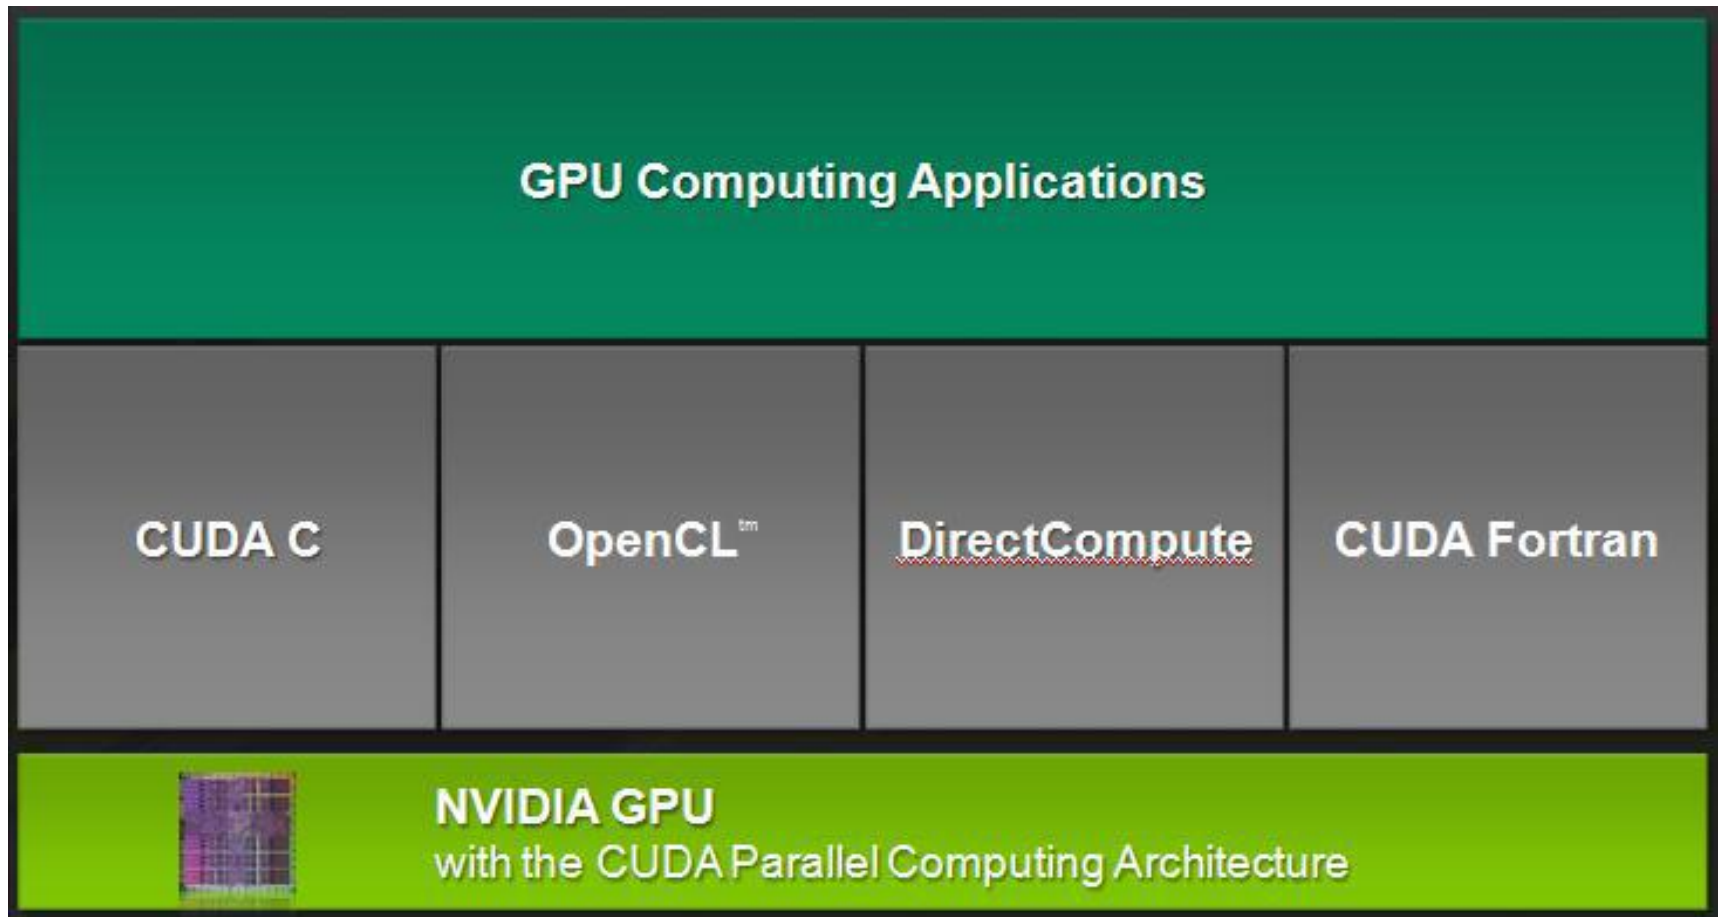


image source: NVIDIA

- declspecs
 - global
 - device
 - shared
 - local
 - constant
- keywords
 - threadIdx
 - blockIdx
- intrinsics
 - `__syncthreads`

```
__device__ float filter[N];

__global__ void convolve (float *image)
{
    __shared__ float region[M];

    ...

    region[threadIdx.x] = image[i];

    __syncthreads ();

    ...

    image[j] = result;
}
```





- runtime API
 - memory management
 - symbol management
 - execution management
- function launch

```
// Allocate GPU memory
void *myimage = cudaMalloc(bytes);

// 100 blocks
// 10 threads per block
convolve<<<100, 10>>>(myimage);
```



CUDA Programming Model

- the GPU (graphics processing unit) is viewed as a compute **device** that
 - is a coprocessor to the CPU or **host**
 - has its own DRAM (**device memory**)
 - runs many **threads in parallel**
- data-parallel portions of an application are executed on the device as **kernels** running in parallel on many threads
- differences between GPU and CPU threads
 - GPU threads are extremely lightweight
 - very little creation overhead
 - GPU needs 1000s of threads for full efficiency
 - multi-core CPU needs only a few



helloWorld.cu



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
#include <helper_cuda.h> // various definitions, macros, ...
#include "helloWorld_kernel.cu.h" // include the code actually
                                // extended on the GPU

#define BLOCK_SIZE 16 // defining constants
#define GW (5 * BLOCK_SIZE)

int main(int argc, char** argv) // the main program
{
    printf("block size: %u \n", BLOCK_SIZE); // debug output on host
    printf("grid dimension: %u\n", GW / BLOCK_SIZE);

    dim3 threads(BLOCK_SIZE); // number of threads per block
    dim3 grid(GW / BLOCK_SIZE); // number of blocks in grid

    helloWorld<<<grid, threads>>>(); // start kernel (executed on device)

    getLastCudaError("Kernel execution failed"); // check for success
}
```



helloWorld_kernel.cu.h



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
#pragma once // avoid multiple reads and
              // interpretation of file

#include <stdio.h>

__global__ void helloWorld() // definition of a global fn
{ // callable from host
    if (blockIdx.x == 0 && threadIdx.x == 0) { // executed only once
        printf("Welcome to the Hello World test " // welcome message
               "program for CUDA.\n");
    }

    // output per thread
    printf("Hello World!\n");
    printf("Block index:  x: %u\n",  blockIdx.x);
    printf("Thread index: x: %u\n",  threadIdx.x);
}
```



Hello (CUDA) World! - Output

```
block size: 16
grid dimension: 5
Welcome to the Hello World test program for CUDA.
Hello World!
Block index:    x: 0
Thread index:   x: 0
Hello World!
Block index:    x: 0
Thread index:   x: 1
Hello World!
Block index:    x: 0
Thread index:   x: 2

[...]

Block index:    x: 4
Thread index:   x: 14
Hello World!
Block index:    x: 4
Thread index:   x: 15
```



Blocks and Grids of Threads

- thread batching
- a kernel is executed as a **grid of thread blocks**
 - all threads share data memory space (**global memory**)
- a **thread block** is a batch of threads that can **cooperate**
 - synchronizing their execution
 - efficiently sharing data through **shared memory**
- two threads from two different blocks cannot cooperate




image source: NVIDIA



Blocks and Grids of Threads

- corresponding code in helloWorld.cu

```
#define BLOCK_SIZE 16  
#define GW (5 * BLOCK_SIZE)  
  
dim3 threads(BLOCK_SIZE);  
dim3 grid(GW / BLOCK_SIZE);  
helloWorld<<<grid, threads>>>();
```

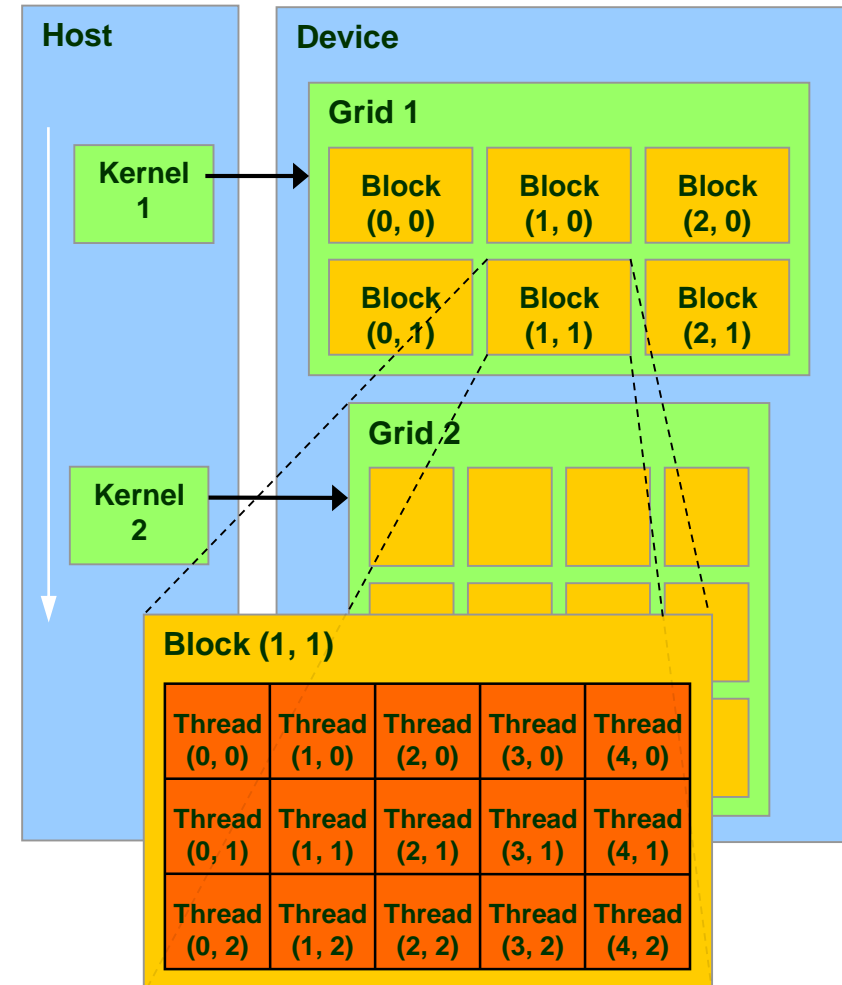


image source: NVIDIA

Block and Thread IDs

- threads and blocks have IDs
 - so each thread can decide what data to work on
 - block ID: 1D, 2D or 3D (>SM2.0)
 - thread ID: 1D, 2D, or 3D
 - data type dim3
- simplifies memory addressing when processing multidimensional data
 - image processing
 - solving PDEs on volumes
 - ...

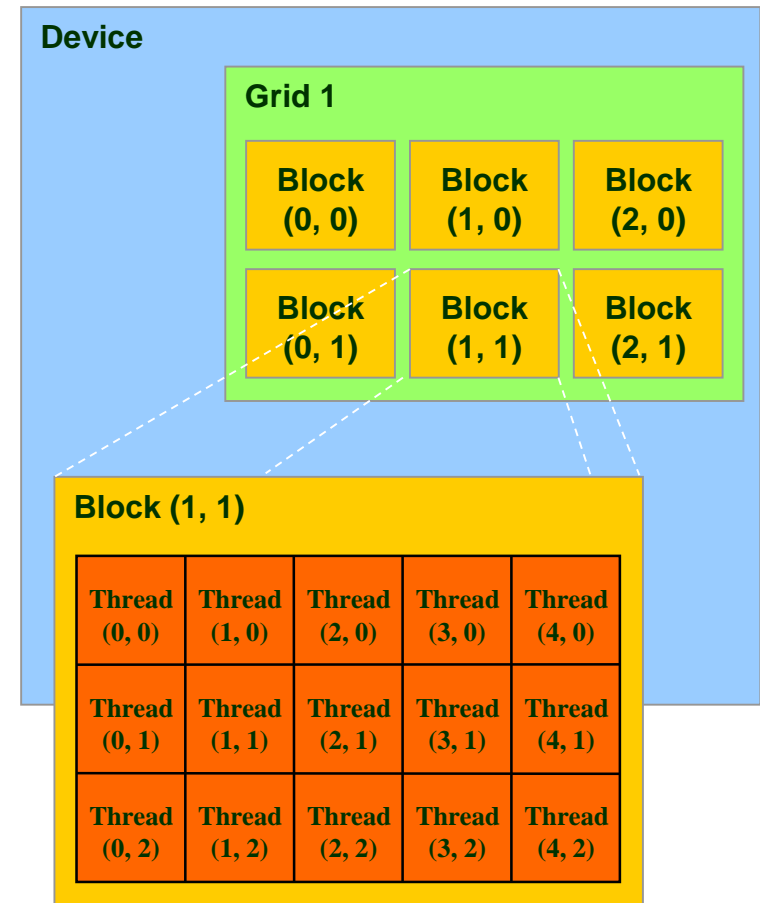


image source: NVIDIA



Block and Thread IDs

- corresponding code in `helloWorld_kernel.cu.h`

```
if (blockIdx.x == 0 && threadIdx.x == 0) {  
    printf("Welcome to the Hello World test "  
        "program for CUDA.\n");  
}  
  
// output per thread  
printf("Hello World!\n");  
printf("Block index:   x: %u\n",   blockIdx.x);  
printf("Thread index:  x: %u\n",   threadIdx.x);
```

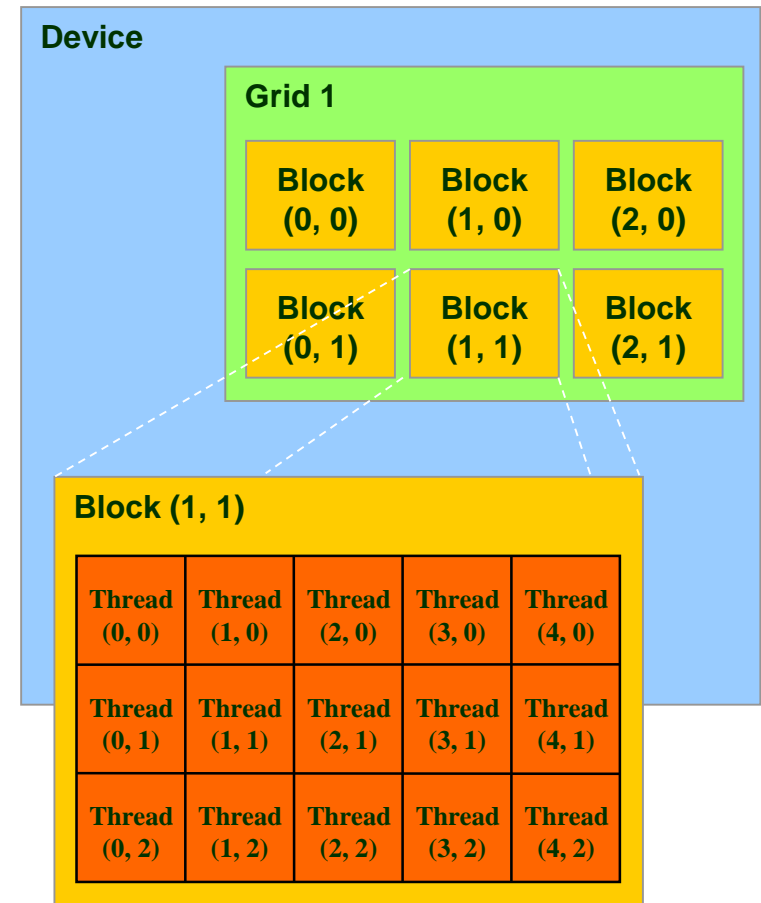


image source: NVIDIA



CUDA Device Memory Space Overview

- Each thread can:
 - R/W per-thread **registers**
 - R/W per-thread **local memory**
 - R/W per-block **shared memory**
 - R/W per-grid **global memory**
 - Read only per-grid **constant memory**
 - Read only* per-grid **texture memory**
- The host can
 - R/W **global memory**
 - R/W **constant memory**
 - R/W **texture memory**

* Can use **surface memory** for R/W access

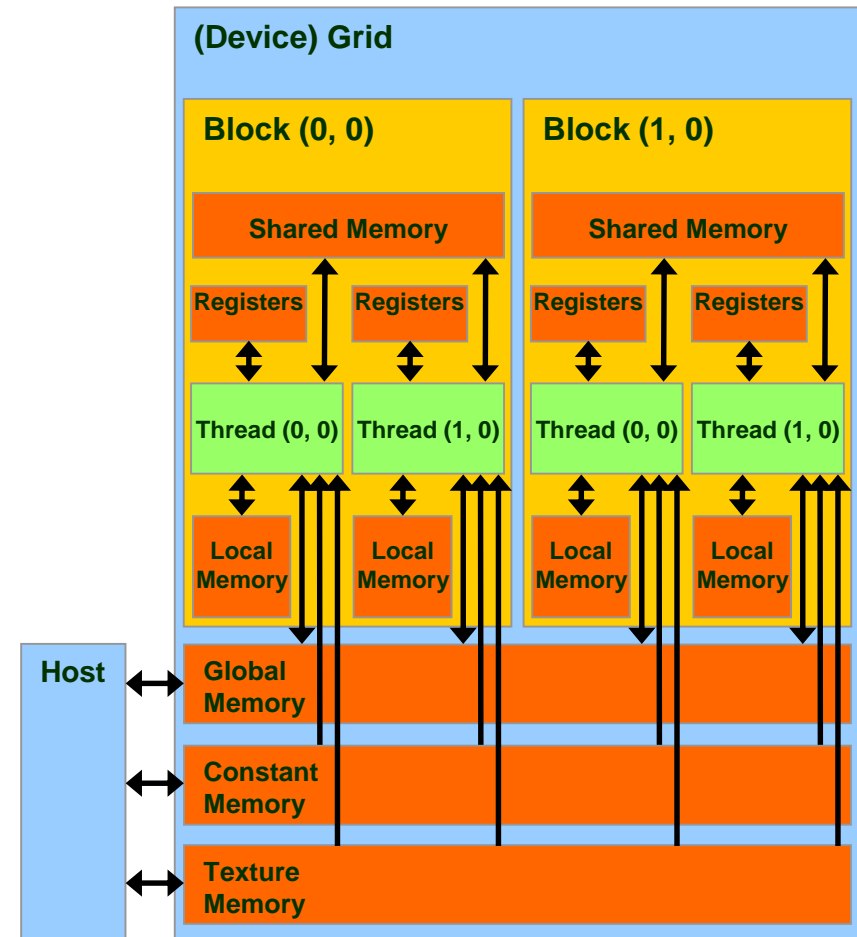


image source: NVIDIA

Global, Constant, and Texture Memories

- long latency access
 - 400-600 clock cycles
- global memory
 - main means of communicating R/W data between **host** and **device**
 - contents visible to all threads
- texture and constant memories
 - constants initialized by host
 - contents visible to all threads
 - cached access
 - texture has built-in interpolation

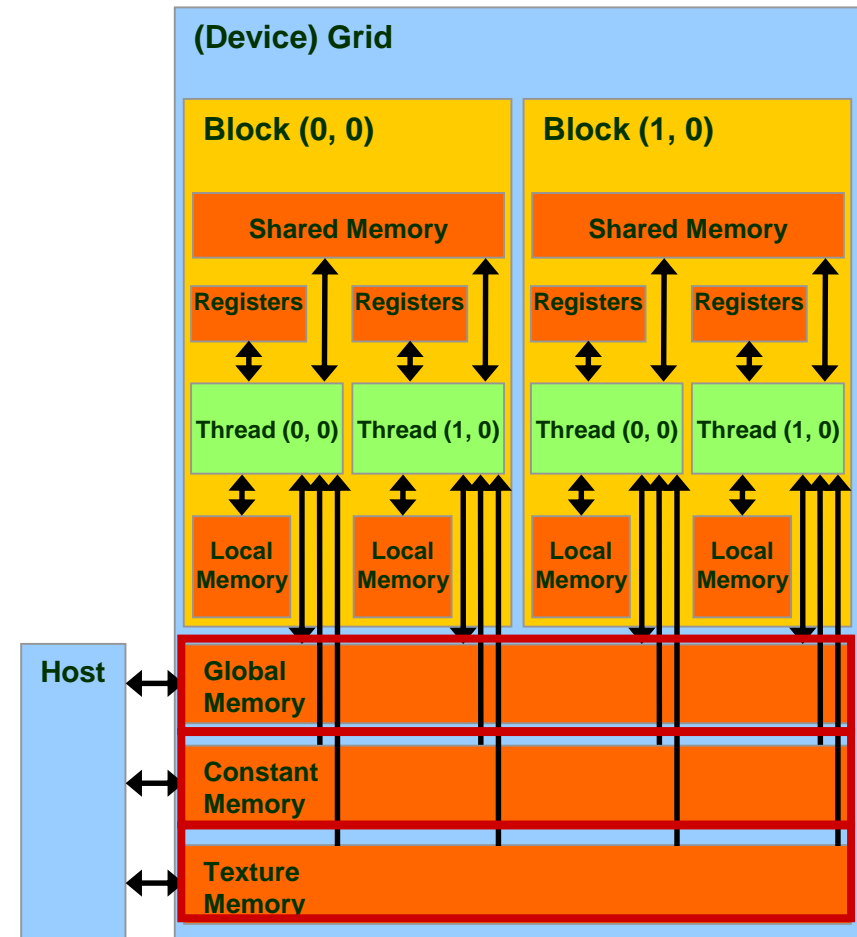


image source: NVIDIA

The CUDA API

- the API is an **extension to the C/C++ programming language**
 - ➔ low learning curve
- the hardware is **designed to enable a lightweight runtime environment and driver**
 - ➔ high performance



A Small Detour: A Matrix Data Type

- NOT part of CUDA
- It will be frequently used in many code examples
 - 2D matrix
 - single precision float elements
 - width * height elements
 - pitch meaningful for alignment and if the matrix is actually a sub-matrix of another matrix
 - data elements allocated and attached to elements

```
typedef struct {  
    int width;  
    int height;  
    int pitch;  
    float* elements;  
} Matrix;
```



CUDA Device Memory Allocation

- `cudaMalloc()`
 - allocates object in the device global memory
 - requires two parameters
 - **address of a pointer** to the allocated object
 - **size** of allocated object
- `cudaFree()`
 - frees object from device global memory
 - pointer to freed object

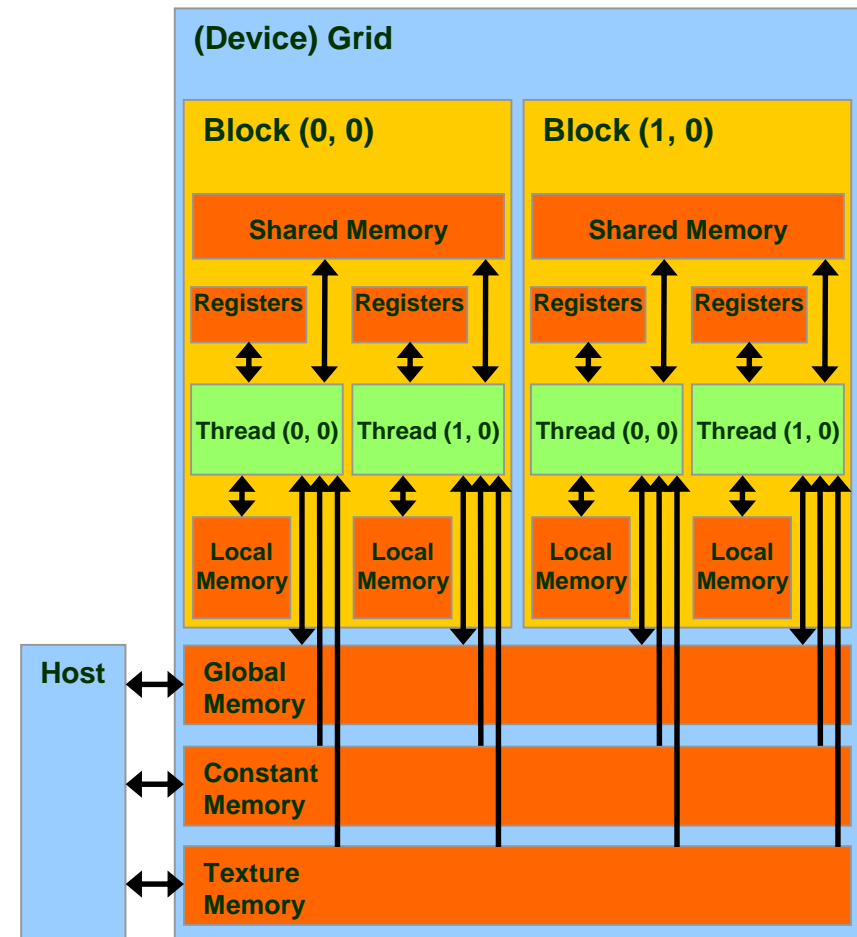


image source: NVIDIA

- code example:
 - allocate a 64x64 float array
 - attach the allocated storage to `Md.elements`
 - “d” is often used to indicate a device data structure

```
#define MATRIX_WIDTH 64
Matrix Md;
int size =
    MATRIX_WIDTH * MATRIX_WIDTH * sizeof(float);
cudaMalloc((void**) &Md.elements, size);
cudaFree(Md.elements);
```



CUDA Host-Device Data Transfer

- `cudaMemcpy()`
 - memory data transfer
 - requires four parameters
 - pointer to destination
 - pointer to source
 - number of bytes copied
 - type of transfer
 - Host to Host
 - Host to Device
 - Device to Host
 - Device to Device
- asynchronous since CUDA 1.1
 - `cudaMemcpyAsync()`

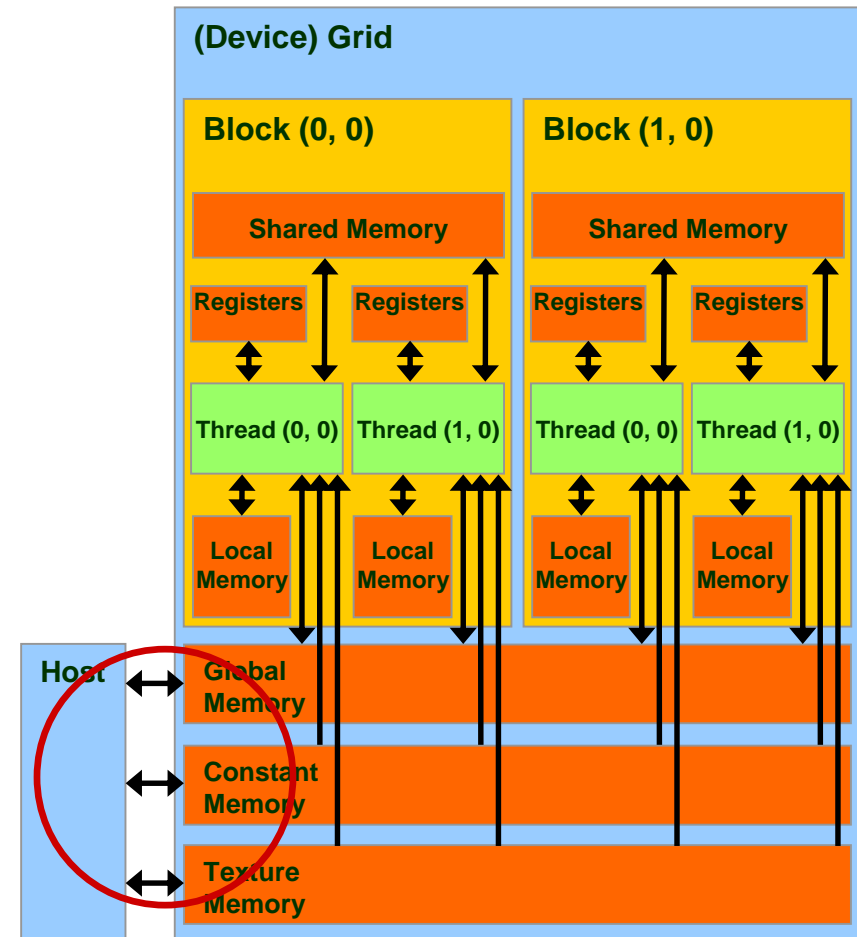


image source: NVIDIA

- code example:
 - transfer a $64 * 64$ single precision float array
 - M is in host memory and Md is in device memory
 - `cudaMemcpyHostToDevice` and `cudaMemcpyDeviceToHost` are symbolic constants

```
cudaMemcpy(Md.elements, M.elements, size,  
           cudaMemcpyHostToDevice);
```

```
cudaMemcpy(M.elements, Md.elements, size,  
           cudaMemcpyDeviceToHost);
```

CUDA Function Declarations

	Executed on	Only callable from
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

- `__global__` defines a kernel function
 - must return `void`
 - can access `blockIdx` and `threadIdx`
- `__device__` and `__host__` can be used together
 - two versions of the function will then be compiled: one for the host and one for the device



- `__device__` functions cannot have their address taken
- for functions executed on the device:
 - recursion only on cards with compute capability $\geq 2.x$
 - no static variable declarations inside the function
 - no variable number of arguments

Calling a Kernel Function

- thread creation
- kernel function must be called with an **execution configuration**

```
__global__ void kernelFunc(...);  
dim3    dimGrid(100, 50);    // 5000 thread blocks  
dim3    dimBlock(4, 8, 8);   // 256 threads per block  
size_t  sharedMemBytes = 64; // 64 bytes of shared memory  
kernelFunc<<<dimGrid, dimBlock, sharedMemBytes>>>(...);
```

- any call to a kernel function is asynchronous, explicit synch needed for blocking



Simple Example: Matrix Multiplication

- straightforward matrix multiplication example
- illustrates the basic features of memory and thread management in CUDA programs
 - leave shared memory usage until later
 - local, register usage
 - thread ID usage
 - memory data transfer API between host and device



Square Matrix Multiplication Example

- $P = M * N$ of size WIDTH x WIDTH
- without tiling
 - one **thread** handles one element of P
 - **M and N are loaded WIDTH times** from global memory

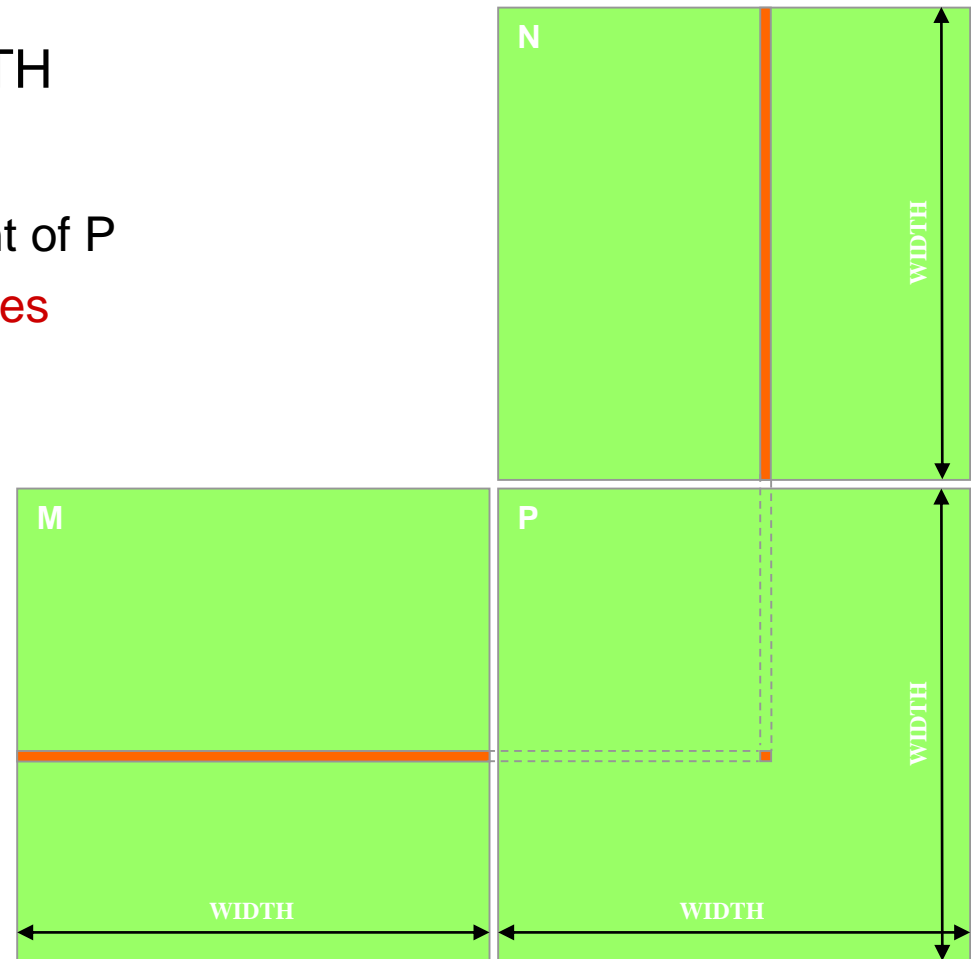


image source: NVIDIA

Step 1: Matrix Data Transfers



```
// Allocate the device memory where we will copy M to
Matrix Md;
Md.width  = WIDTH;
Md.height = WIDTH;
Md.pitch  = WIDTH;

int size = WIDTH * WIDTH * sizeof(float);
cudaMalloc((void**)&Md.elements, size);

// Copy M from the host to the device
cudaMemcpy(Md.elements, M.elements, size,
            cudaMemcpyHostToDevice);

...

// Read M from the device to the host into P
cudaMemcpy(P.elements, Md.elements, size,
            cudaMemcpyDeviceToHost);

// Free device memory
cudaFree(Md.elements);
```



Step 2: Matrix Multiplication

- simple host code in C
 - multiplication on the CPU host in double precision
 - for simplicity, we will assume that all dimensions are equal

```
void MatrixMulOnHost(const Matrix M, const Matrix N, Matrix P)
{
    for (int i = 0; i < M.height; ++i) {
        for (int j = 0; j < N.width; ++j) {
            double sum = 0;
            for (int k = 0; k < M.width; ++k) {
                double a = M.elements[i * M.width + k];
                double b = N.elements[k * N.width + j];
                sum += a * b;
            }
            P.elements[i * N.width + j] = sum;
        }
    }
}
```



Multiply Using One Thread Block

- one block of threads compute matrix P
 - one element of P per thread
- each thread
 - loads a row of matrix M
 - loads a column of matrix N
 - perform one multiply and addition per pair of M and N elements
 - compute to off-chip memory access ratio close to 1:1 (not very high)
- size of matrix limited by the number of threads allowed in a thread block




image source: NVIDIA

Step 3: Matrix Mul. – Host-Side Code 1



```
int main(void)
{
    // Allocate and initialize the matrices
    Matrix M = AllocateMatrix(WIDTH, WIDTH, 1);
    Matrix N = AllocateMatrix(WIDTH, WIDTH, 1);
    Matrix P = AllocateMatrix(WIDTH, WIDTH, 0);

    // M * N on the device
    MatrixMulOnDevice(M, N, P);

    // Free matrices
    FreeMatrix(M);
    FreeMatrix(N);
    FreeMatrix(P);

    return EXIT_SUCCESS;
}
```



Step 3: Matrix Mul. – Host-Side Code 2



```
// Matrix multiplication on the device
void MatrixMulOnDevice(const Matrix M, const Matrix N, Matrix P)
{
    // Load M and N to the device
    Matrix Md = AllocateDeviceMatrix(M);
    CopyToDeviceMatrix(Md, M);
    Matrix Nd = AllocateDeviceMatrix(N);
    CopyToDeviceMatrix(Nd, N);

    // Allocate P on the device
    Matrix Pd = AllocateDeviceMatrix(P);
    CopyToDeviceMatrix(Pd, P); // Clear memory
}
```



Step 3: Matrix Mul. – Host-Side Code 3



```
// Setup the execution configuration
dim3 dimBlock(WIDTH, WIDTH);
dim3 dimGrid(1, 1);

// Launch the device computation threads!
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd);

// Read P from the device
CopyFromDeviceMatrix(P, Pd);

// Free device matrices
FreeDeviceMatrix(Md);
FreeDeviceMatrix(Nd);
FreeDeviceMatrix(Pd);
}
```



Step 4: Device-side Kernel Function 1



```
// Matrix multiplication kernel - thread specification
__global__ void MatrixMulKernel(Matrix M, Matrix N, Matrix P) {
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // pValue is used to store the element of the matrix
    // that is computed by the thread
    double pValue = 0;
```



Step 4: Device-side Kernel Function 2



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
for (int k = 0; k < M.width; ++k) {  
    float Melement =  
        M.elements[ty * M.pitch + k];  
    float Nelement =  
        N.elements[k * N.pitch + tx];  
    pValue += Melement * Nelement;  
}  
// Write the matrix  
// to device memory;  
// each thread writes  
// one element  
P.elements[ty * P.pitch  
    + tx] = pValue;  
}
```

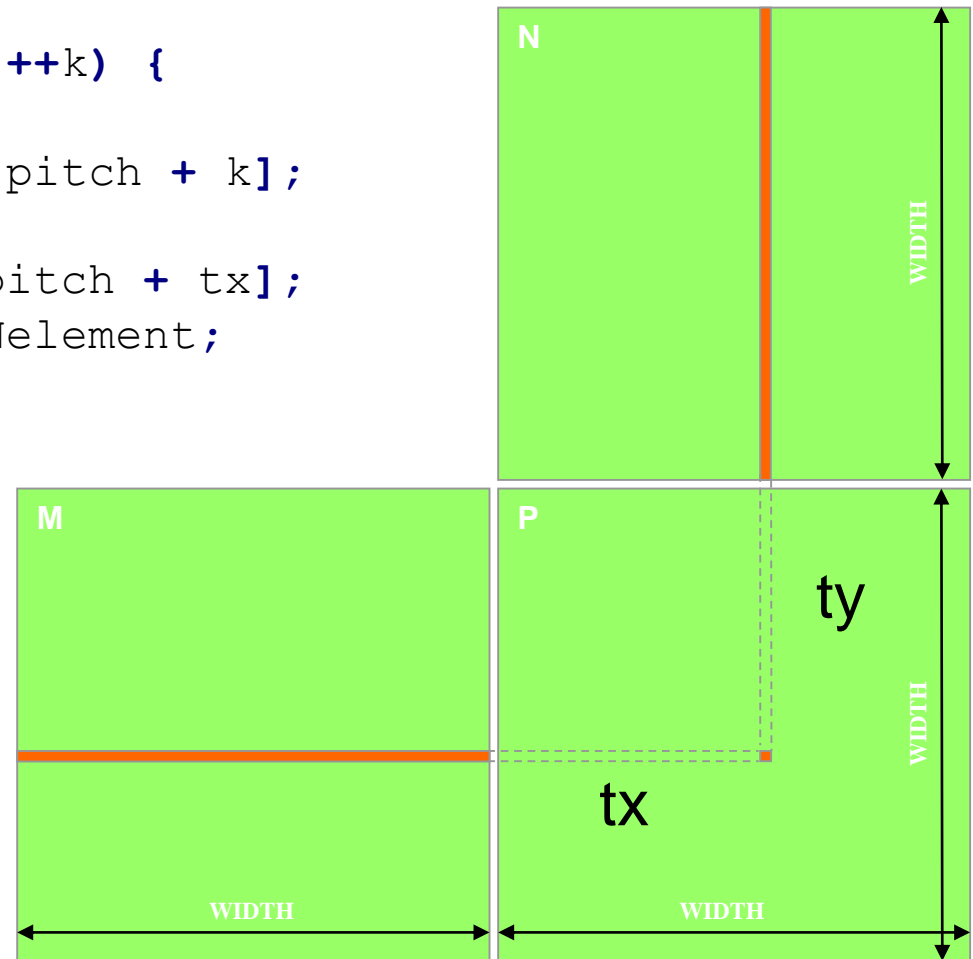


image source: NVIDIA



Step 5: Some Loose Ends 1



```
// Allocate a device matrix of same size as M.
Matrix AllocateDeviceMatrix(const Matrix M)
{
    Matrix Mdevice = M;
    int size = M.width * M.height * sizeof(float);
    cudaMalloc((void**) &Mdevice.elements, size);
    return Mdevice;
}

// Free a device matrix.
void FreeDeviceMatrix(Matrix M)
{
    cudaFree(M.elements);
}

void FreeMatrix(Matrix M)
{
    free(M.elements);
}
```



Step 5: Some Loose Ends 2

```
// Copy a host matrix to a device matrix.
void CopyToDeviceMatrix(Matrix Mdevice, const Matrix Mhost)
{
    int size = Mhost.width * Mhost.height * sizeof(float);
    cudaMemcpy(Mdevice.elements, Mhost.elements, size,
               cudaMemcpyHostToDevice);
}

// Copy a device matrix to a host matrix.
void CopyFromDeviceMatrix(Matrix Mhost, const Matrix Mdevice)
{
    int size=Mdevice.width * Mdevice.height * sizeof(float);
    cudaMemcpy(Mhost.elements, Mdevice.elements, size,
               cudaMemcpyDeviceToHost);
}
```



Step 6: Arbitrary Sized Square Matrices

- each 2D thread block computes a $(\text{BLOCK_WIDTH})^2$ sub-matrix (tile) of the result matrix
 - each has $(\text{BLOCK_WIDTH})^2$ threads
- 2D grid
 - $(\text{WIDTH}/\text{BLOCK_WIDTH})^2$ blocks
- need to put a loop around the kernel call for cases where WIDTH is greater than max grid size!

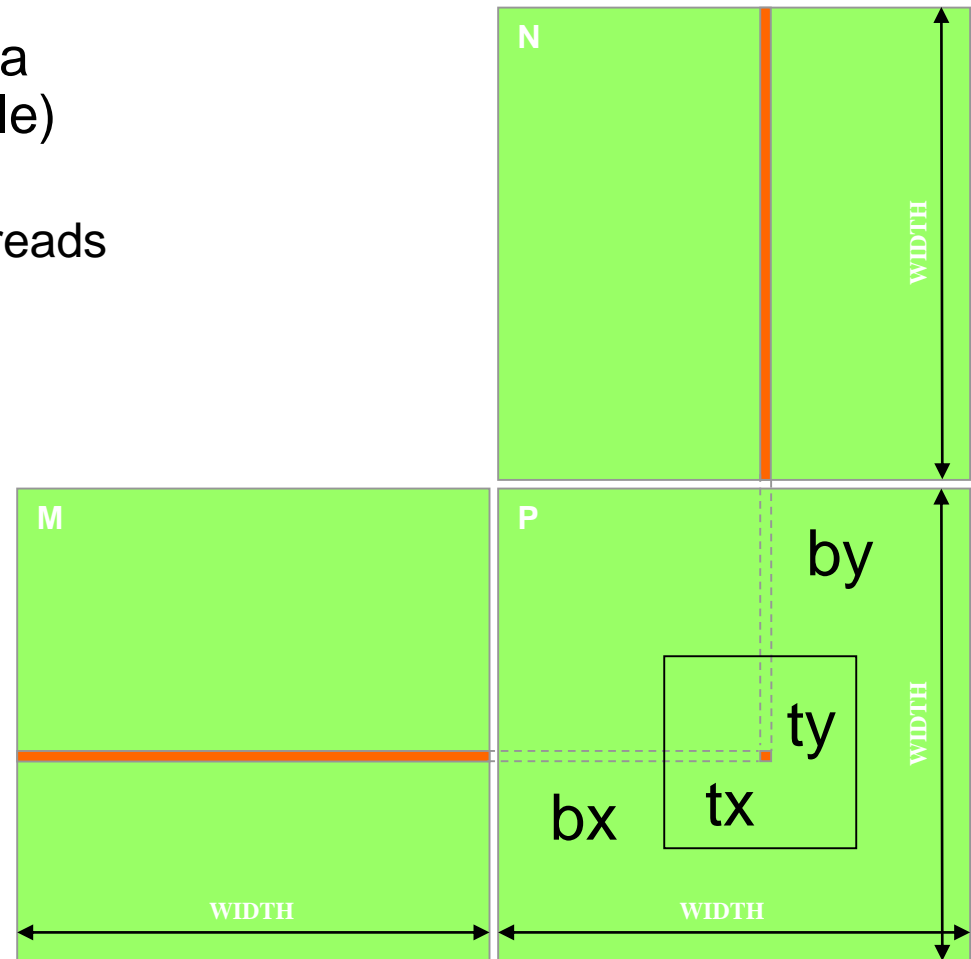


image source: NVIDIA

Assignments and Accounts

- Exercise 1 – due Monday 2.11.2015
- read Chapters 1+2 of the CUDA 7.5 Programming Guide
 - available on NVIDIA's web page <http://docs.nvidia.com/>
- accounts will be handed out to you ASAP

