

**PMPP 2015/16**



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

# Design Patterns (2)



# Course Schedule


---

12.10.2015	Introduction to PMPP
13.10.2015	Lecture CUDA Programming 1
19.10.2015	Lecture CUDA Programming 2
20.10.2015	Lecture CUDA Programming 3
26.10.2015	Lecture Parallel Basics, <a href="#">Exercise 1 assigned</a>
27.10.2015	Questions and Answers (Q&A), S3 19, Room 2.8
2.11.2015	<a href="#">Intro Final Proj.</a> , <a href="#">Ex. 1 due</a> , <a href="#">Ex. 2 assigned</a> , Lecture PRAM
3.11.2015	Lecture PRAM (2)
9.11.2015	<a href="#">Final Projects assigned</a> , L. Parallel Sort., <a href="#">Exercise 2 due</a>
10.11.2015	Questions and Answers (Q&A)
16.11.2015	Questions and Answers (Q&A)
17.11.2015	Questions and Answers (Q&A)
23.11.2015	<a href="#">1<sup>st</sup> Status Presentation Final Projects</a>
24.11.2015	<a href="#">1<sup>st</sup> Status Presentation Final Projects (continued)</a>
30.11.2015	Lecture Design Patterns
1.12.2015	Questions and Answers (Q&A)

---



# (Preliminary) Course Schedule

 you are here	7.12.2015	Lecture Design Patterns (2), Performance Tuning
	8.12.2015	Questions and Answers (Q&A)
	14.12.2015	Lecture
	15.12.2015	Questions and Answers (Q&A)
	11.1.2016	2 <sup>nd</sup> Status Presentation Final Projects
	12.1.2016	2 <sup>nd</sup> Status Presentation Final Projects (continued)
	18.1.2016	
	19.1.2016	
	25.1.2016	
	26.1.2016	
	1.2.2016	
	2.2.2016	
	8.2.2016	Final Presentation Final Projects
	9.2.2016	Final Presentation Final Projects (continued)



# Final Projects – Second Presentation

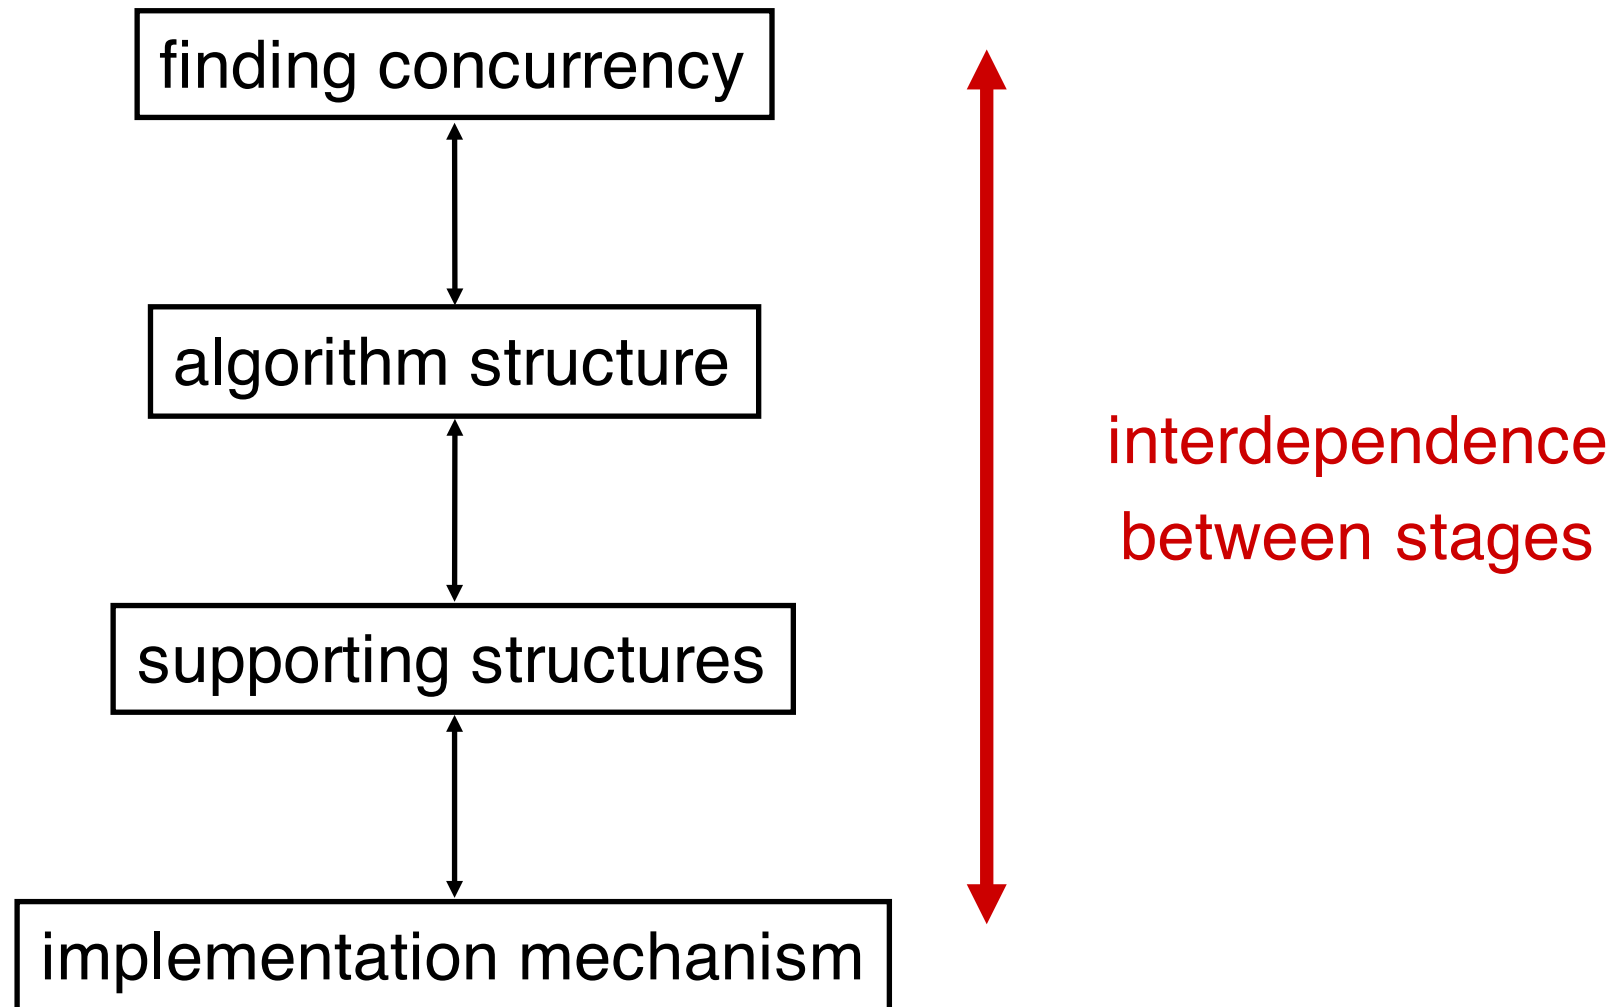
- second presentation on 11.01.2016 and 12.01.2016
  - present the current status of your project
  - discuss planned next steps
  - discuss open issues
  - DO NOT present the topic and the approach you plan to use once more
  - 7 minutes per group (!)
- slides must be submitted no later than 10 am on 11.01.2016
  - one presentation per group (not per topic)
- everybody should give a portion of the presentation
- mandatory (talk to us if this is a problem)

# Written Exam

- Date for the exam will be 02.03.2016, starting at 2pm
  - do not forget to register in TuCAN!
- Rooms: S101/A1 and S101/A01
- The exam will be in English only
- overlap with exam “Communication Networks 2” will be resolved by an addition exam slot for CN2
  - exceptional, one-time solution only
  - thanks to Prof. Steinmetz and his team




# Design Spaces



from Patterns for Parallel Programming by Mattson, Sanders, and Massingill


# Design Spaces and Design Patterns

- finding concurrency
    - programmer working in the problem domain to identify available concurrency and expose it in the algorithm design
- 
- ➔ task decomposition pattern
  - ➔ data decomposition pattern
  - ➔ group tasks pattern
  - ➔ order tasks pattern
  - ➔ data sharing pattern
  - ➔ design evaluation pattern



# Design Spaces and Design Patterns


- algorithm structure
    - programmer working with high-level structures for organizing a parallel algorithm
- 
- ➔ task parallelism pattern
  - ➔ divide and conquer pattern
  - ➔ geometric decomposition pattern
  - ➔ recursive data pattern
  - ➔ pipeline pattern
  - ➔ event-based coordination pattern






# Design Spaces and Design Patterns

- supporting structures
  - shift from algorithms to source code
  - organization of parallel program
  - techniques to manage shared data
- ➔ SPMD pattern
- ➔ master/worker pattern
- ➔ loop parallelism pattern
- ➔ fork/join pattern
- ➔ shared data pattern
- ➔ shared queue pattern
- ➔ distributed array pattern
- ➔ ...

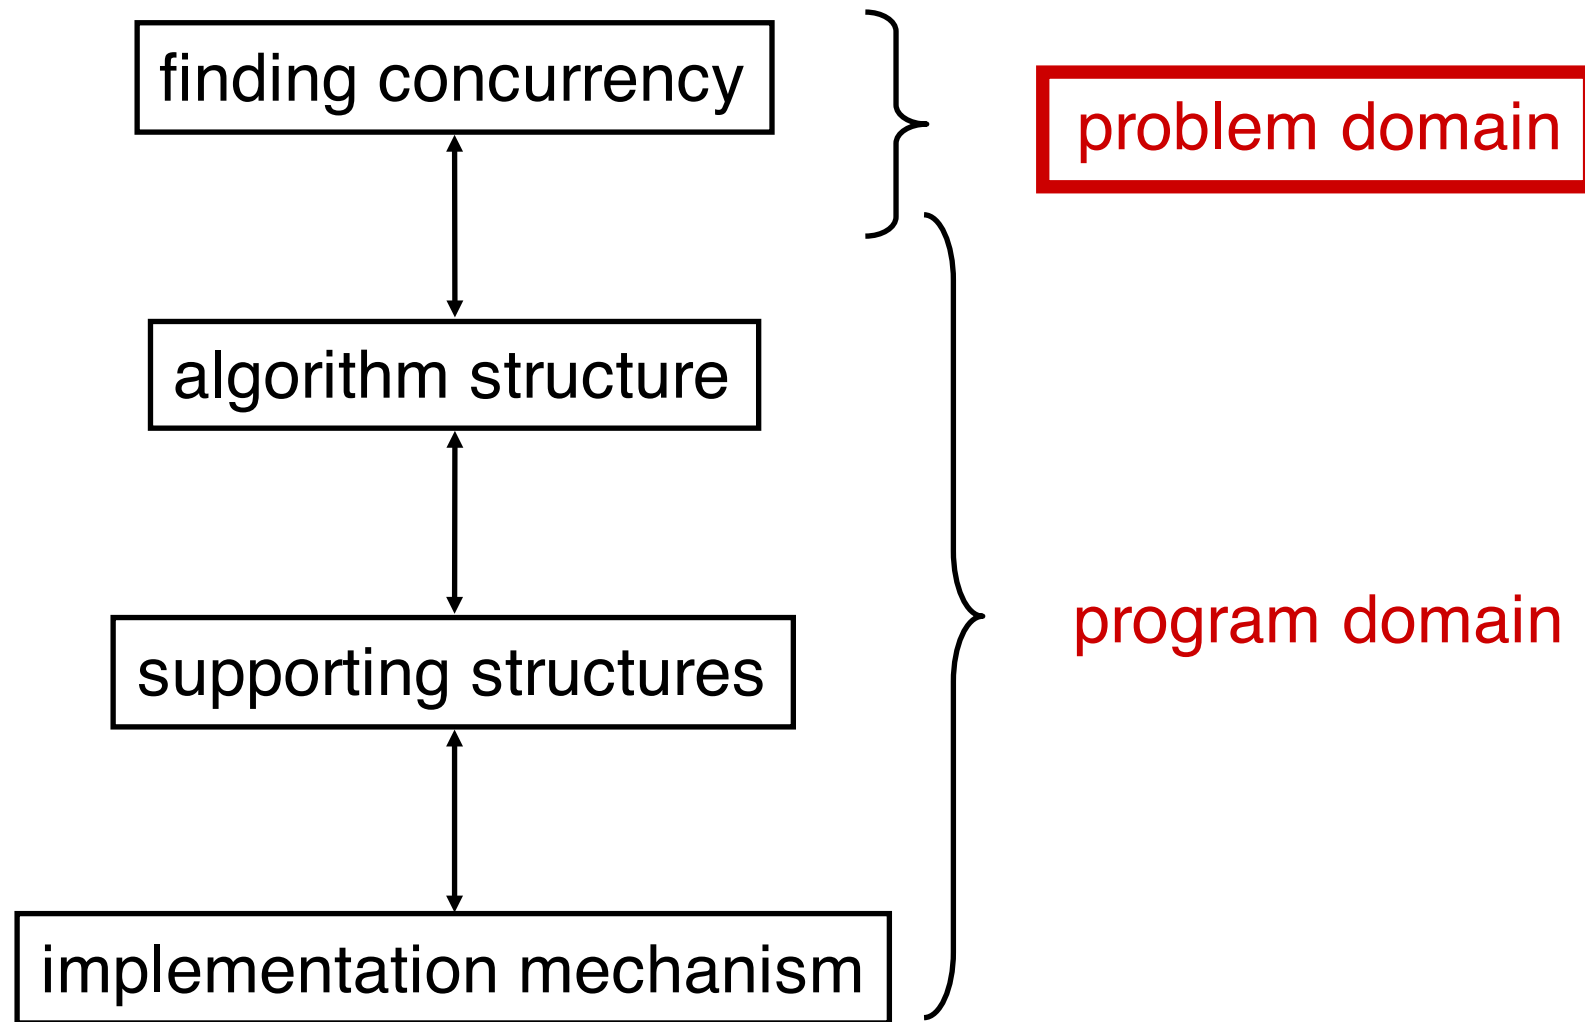


# Design Spaces (and Design Patterns)

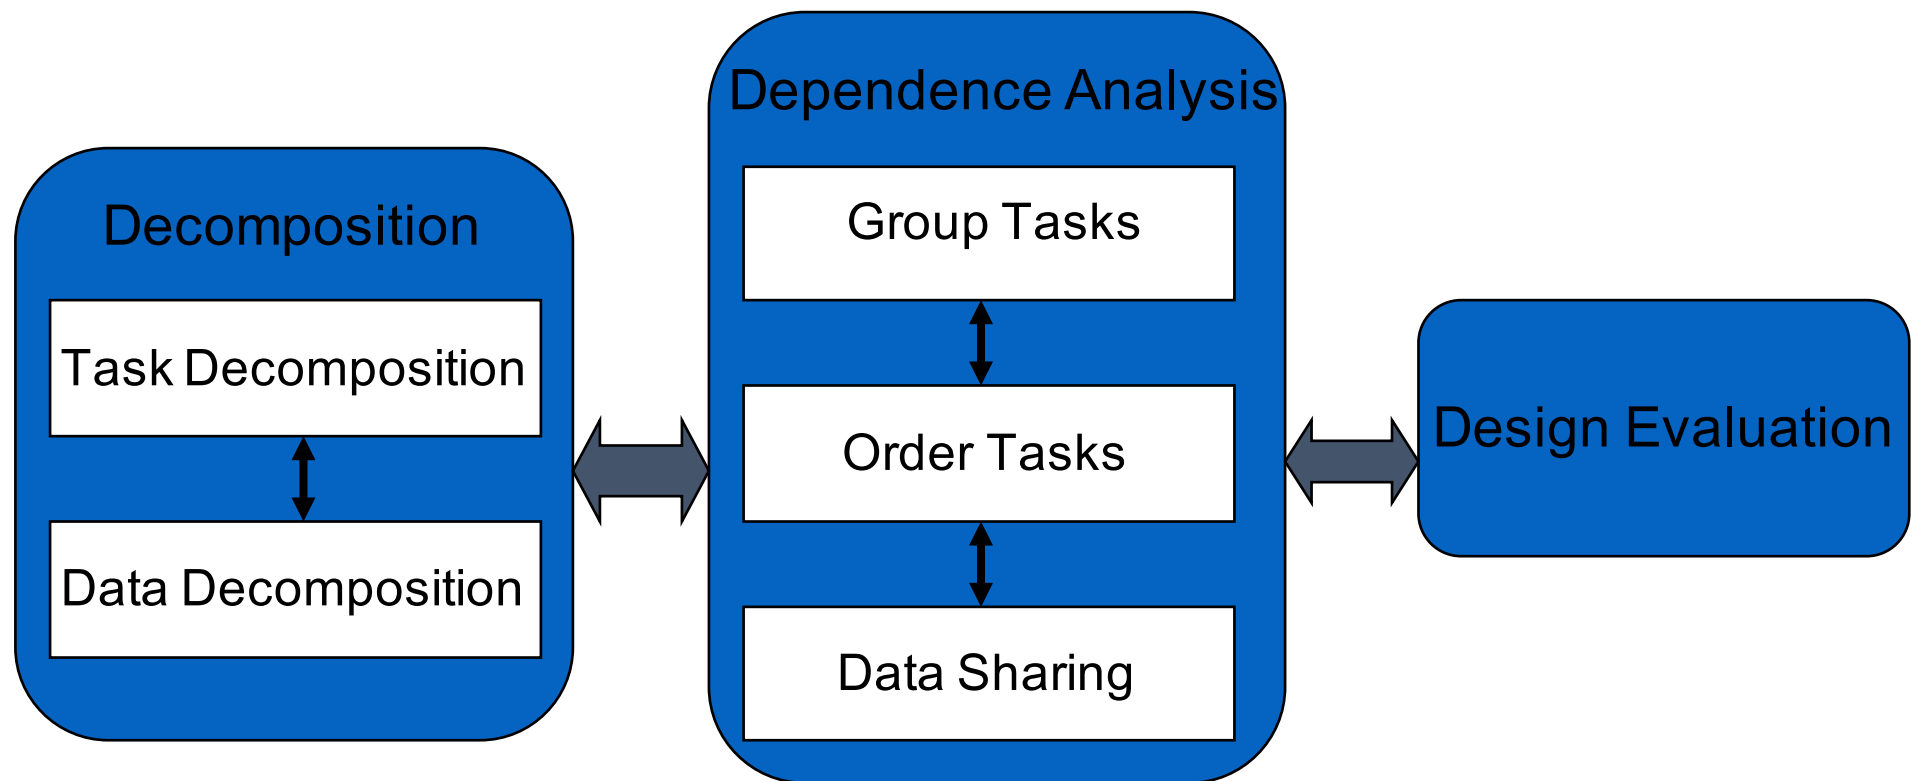
- implementation mechanisms
    - specific software constructs for implementing a parallel program
- UE management
- synchronization
- communication



# Design Spaces

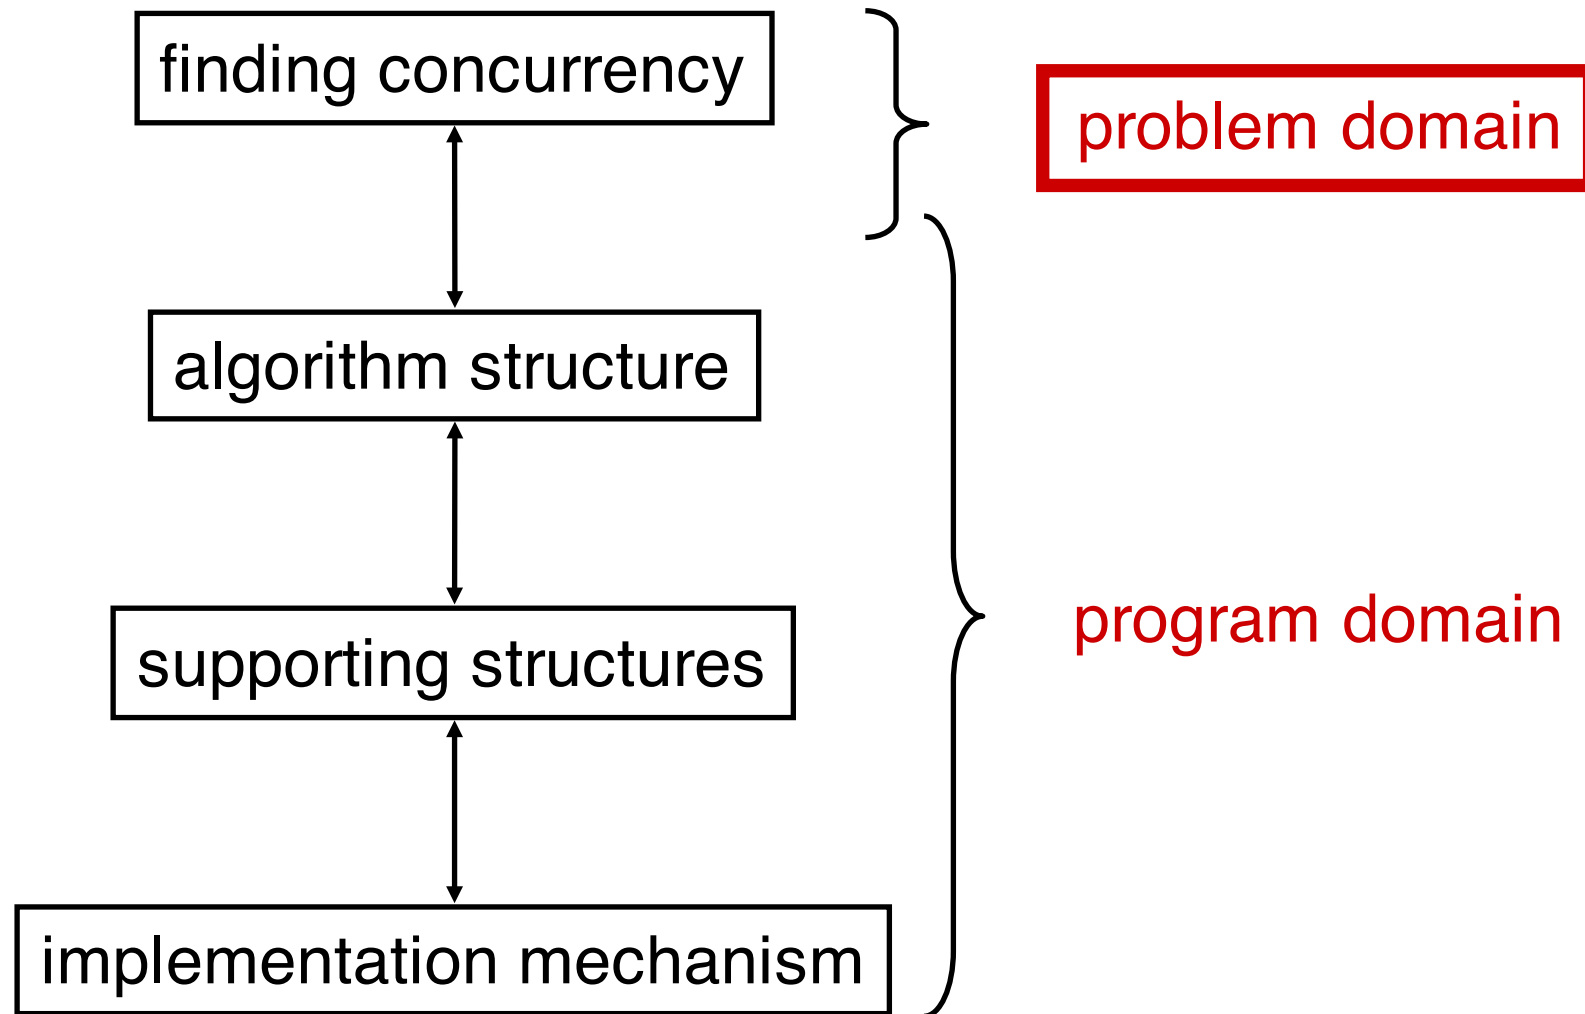


# Finding Concurrency



- typically an iterative process
- opportunities exist for dependence analysis to play earlier role in decomposition

# Design Spaces



# Algorithm

- a step by step procedure that is guaranteed to terminate, such that each step is precisely stated and can be carried out by a computer
  - **definiteness**: the notion that each step is precisely stated
  - **effective computability**: each step can be carried out by a computer
  - **finiteness**: the procedure terminates
- multiple algorithms can be used to solve the same problem
  - some require fewer steps and some exhibit more parallelism



# Algorithm Structure

## Organize by Tasks

Task Parallelism

Divide and Conquer

## Organize by Data

Geometric  
Decomposition

Recursive Data  
Decomposition

## Organize by Flow

Pipeline

Event Condition

- important to re-evaluate the design
  - especially suitability for the target platform

# Algorithm Structure – Main Forces

- efficiency
  - program should run quickly and make good use of resources
- simplicity
  - easy to understand code (develop, debug, verify, modify)
- portability
  - should run on a wide range of computers
  - lifetime of program typically longer than lifetime of computer
  - protects investment in software
- scalability
  - effective for a wide range of processing elements (PE)





# Algorithm Structure – Main Forces

- possible conflicts
  - efficiency vs. portability
    - using special hardware features yields efficient but not portable code
  - efficiency vs. simplicity
    - e.g., task parallelism may require complicated scheduling
- ➔ overall goals
- ➔ balance between abstraction, portability and
  - ➔ suitability for a particular target platform



# Algorithm Structure – Considerations

- target platform
  - should ideally not be necessary at that point in the development
  - but required to get efficient code
- order of magnitude of UEs
  - e.g., 10s or 1000s
- cost of sharing information between UEs
  - shared memory?
- programming environment
  - often multiple environments available for the platform


# Algorithm Structure – Considerations

- major organizing principle
  - should emerge from the finding concurrency design space
- organization by tasks
  - only one group of tasks active at a time
  - interaction between tasks is major feature
  - e.g., embarrassingly parallel programs
- organization by data decomposition
  - e.g., update of a large data structure as main feature of the program
- organization by flow of data
  - e.g., continuous or discrete flow of data

➔ or a combination of the above

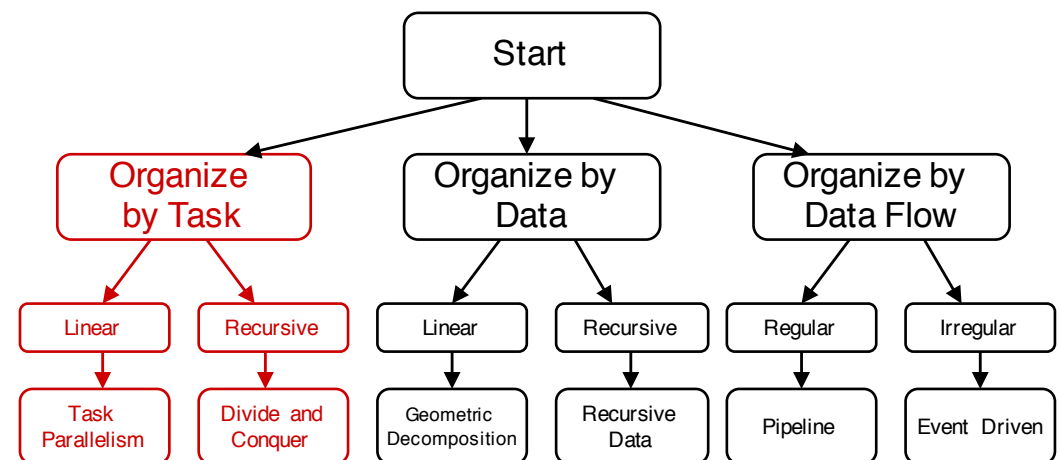


# Choosing Algorithm Structure



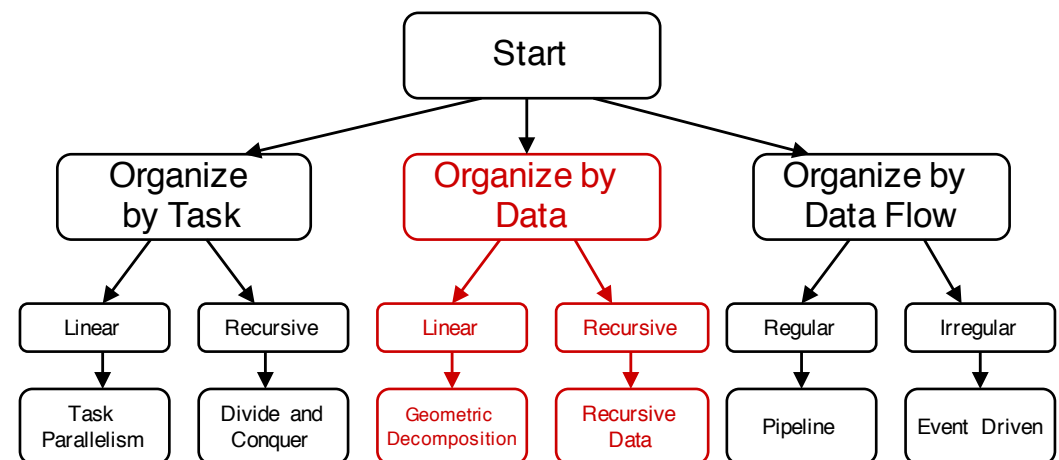
# Choosing Algorithm Structure

- organize by task
  - ➔ **task parallelism**
- execution of tasks best organizing principle
- set of tasks enumerated linear in any number of dimensions
  - ➔ **divide and conquer**



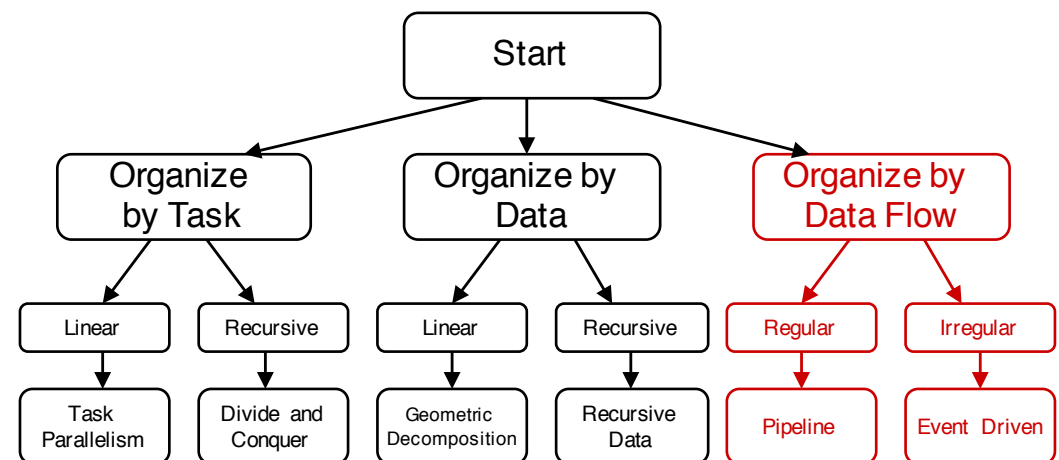
# Choosing Algorithm Structure

- organize by data decomposition
- data decomposition best organizing principle
- program decomposed into discrete subspaces, solutions computed independently interacting with a small number of neighbors
  - **geometric decomposition**
- problem defined as following links in a recursive data structure
  - **recursive data**



# Choosing Algorithm Structure

- organize by flow of data
- flow of data imposes an ordering on groups of tasks
- regular, one-way flow of data which doesn't change during execution
  - ➔ **pipeline**
- irregular, dynamic, not predictable flow of data
  - ➔ **event driven**



# Algorithm Structure

## Organize by Tasks

Task Parallelism

Divide and Conquer

## Organize by Data

Geometric  
Decomposition

Recursive Data  
Decomposition

## Organize by Flow

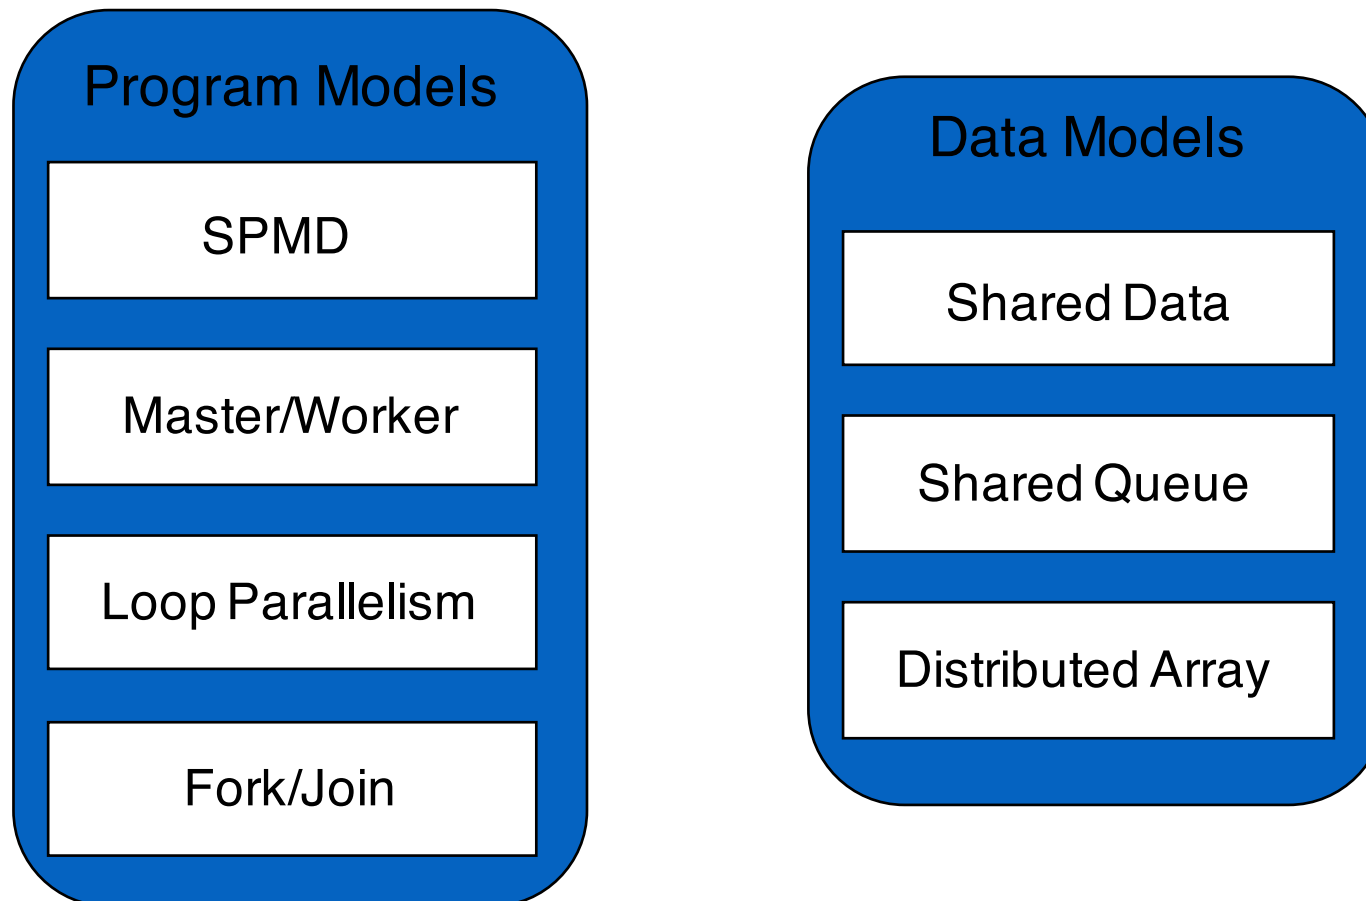
Pipeline

Event Condition

- important to re-evaluate the design
  - especially suitability for the target platform



# Supporting Structures



- supporting structures not necessarily mutually exclusive

# Program Models

- SPMD (Single Program, Multiple Data)
  - all PE' s execute the **same program** in parallel, but each has its **own data**
  - each PE uses a unique ID to access its portion of data
  - different PE can follow different paths through the same code
  - essentially the **CUDA grid model**
  - SIMD/SIMT are special cases, SIMT corresponds to a CUDA warp
- master/worker
  - master creates a **set of worker processes/threads** and a **bag of tasks**
  - tasks are processed by the workers



# Program Models

- loop parallelism
    - runtime of serial program **dominated by a set of compute-intensive loops**
    - **different iterations** of the loop are **executed in parallel**
  - fork/join
    - **main UE forks off other UEs** working in parallel
    - often **forking UE** waits for other UEs to terminate and join
- ➔ master/worker can be implemented using the SPMD or fork/join pattern
- ➔ patterns not exclusive, not unique
- ➔ describe major idioms used by experienced programmers

# Data Structures

- shared data
  - all threads share a major data structure
  - this is what CUDA supports
  - general problem of handling shared data
  - correctness and performance issues
- shared queue
  - all threads see a “thread safe” queue that maintains ordering of data communication
- distributed array
  - decomposed and distributed among threads
  - limited support in CUDA Shared Memory

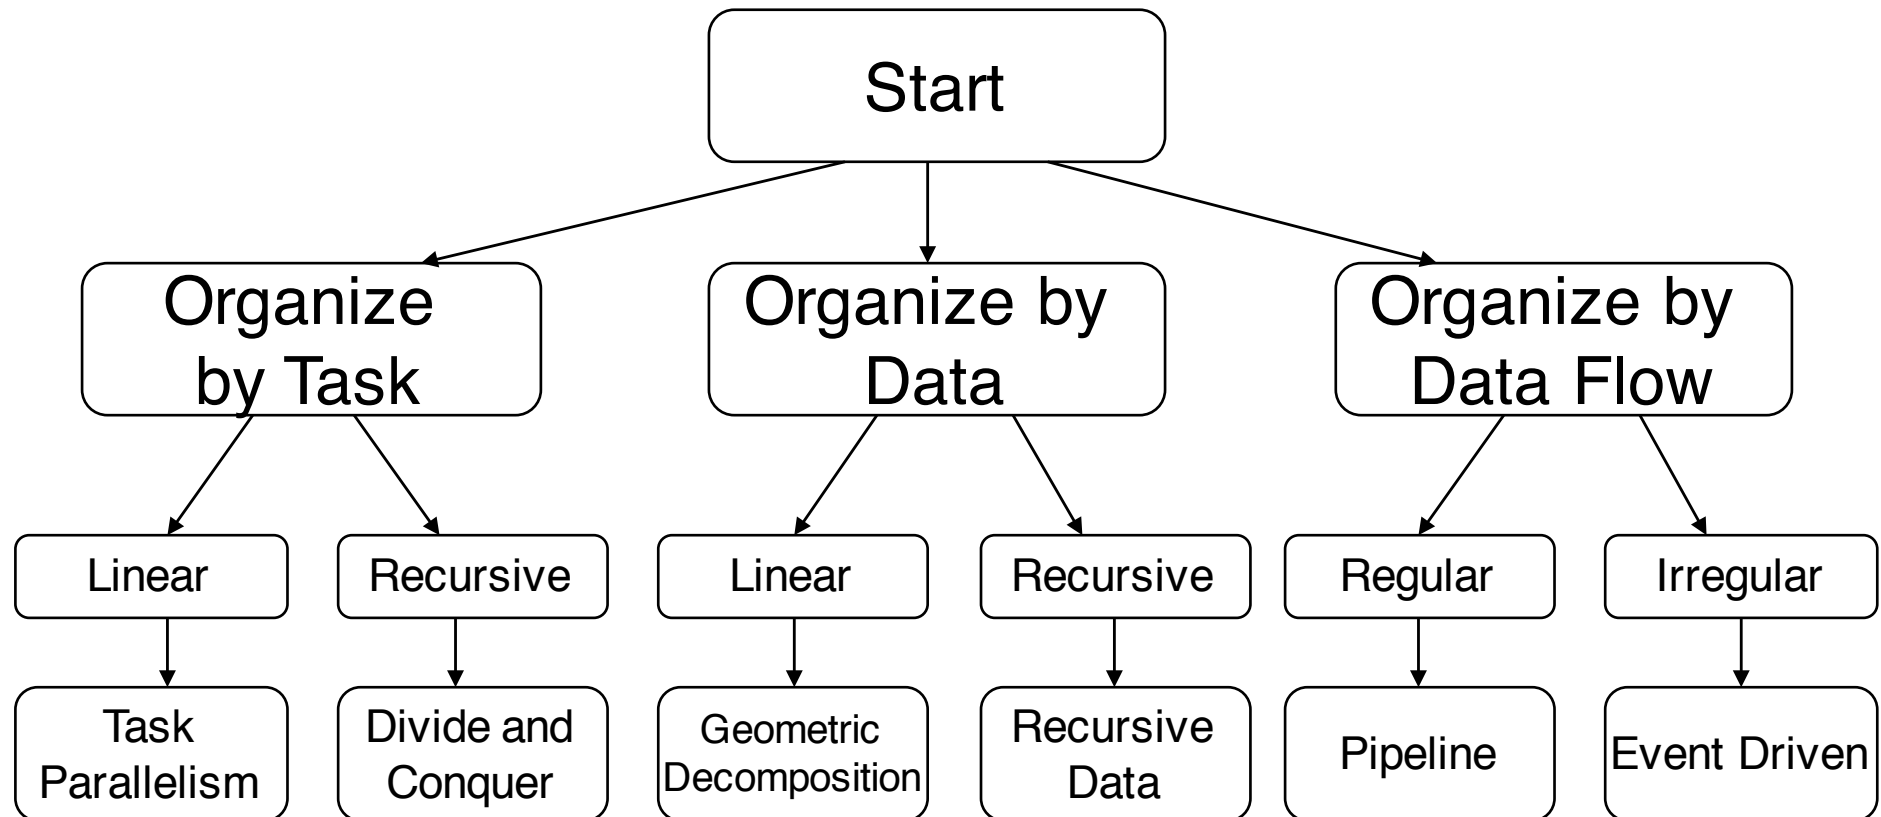


# Supporting Structures: Forces

- clarity of abstraction
  - algorithm clearly apparent from the source code?
- scalability
- efficiency
- maintainability
- environment affinity
  - does the program fit the hardware and programming environment
- sequential equivalence
  - equivalent results to sequential execution if executed on many UEs?
  - relationship sequential/parallel clear?



# Review: Algorithm Structure



# Algorithm Structures vs. Coding Styles

	SPMD	loop parallelism	master/ worker	fork/join
task parallelism	😊😊😊😊	😊😊😊😊	😊😊😊😊	😊😊
divide and conquer	😊😊😊	😊😊	😊😊	😊😊😊😊
geometric decomposition	😊😊😊😊	😊😊😊	😊	😊😊
recursive data	😊😊		😊	
pipeline	😊😊😊		😊	😊😊😊😊
event-based coordination	😊😊		😊	😊😊😊😊


**CUDA**

Source: Mattson et al.



# Design Spaces (and Design Patterns)

- implementation mechanisms
    - specific software constructs for implementing a parallel program
- UE management
- synchronization
- communication





# Implementation Mechanisms

- low level operations unique to parallel programming
  - only small subset used by most programmers
  - UE management
    - how processes and threads are created, managed, destroyed
  - synchronization
    - enforce ordering between events
    - correct access to shared data structures
  - communication
    - information exchange between UEs
- ➔ depends highly on the target platform

# Implementation Mechanisms

- UE management
  - process: heavyweight object with its own state/context
  - thread: lightweight object, part of a process
  - CUDA threads: extremely lightweight
- device threads handled in CUDA as kernel
- host threads in CUDA using, e.g., `cutStartThread`
  - maps to native host threads



# Implementation Mechanisms

- synchronization
- memory synchronization for shared memory access
  - ensure that different threads see the same memory content
  - fences
  - CUDA:
    - volatile shared memory
    - syncthreads command (barrier synchronization)
- barrier synchronization
  - all UEs must arrive at this point before proceeding with computation
  - CUDA: e.g., syncthreads command
- mutual exclusion
  - only one UE can process a critical section
  - CUDA: e.g., atomic operations

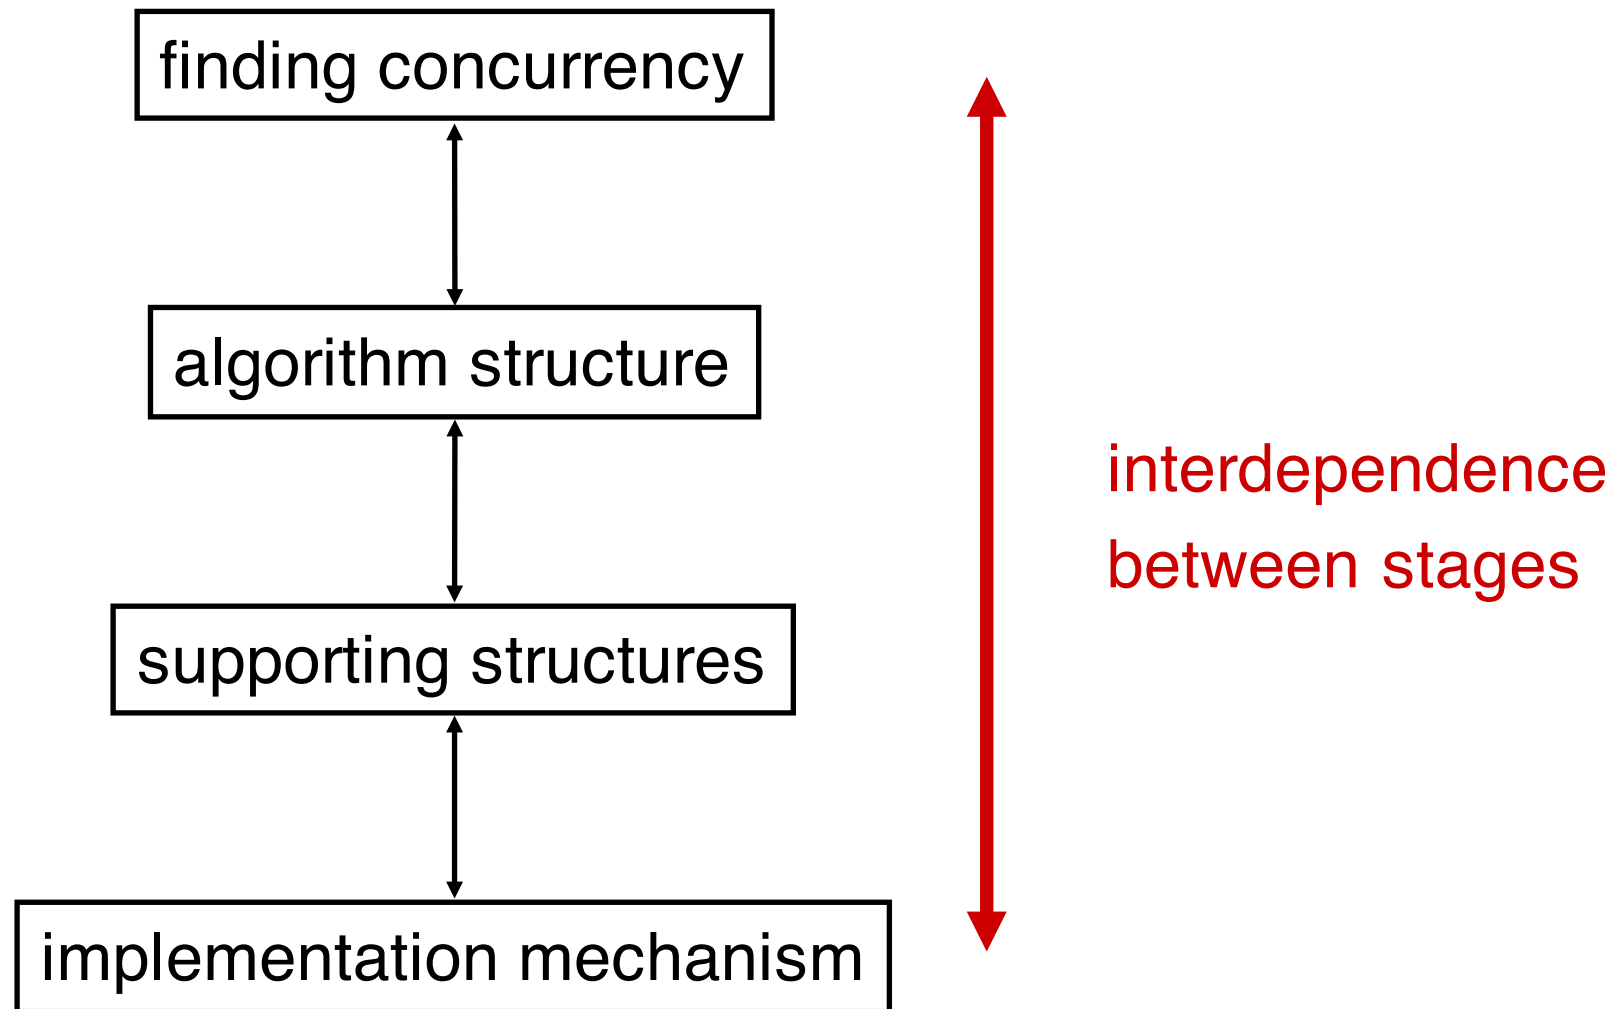


# Implementation Mechanism

- communication
- message passing between pairs of UEs
- collective communication
  - broadcast
    - single message sent to all UEs
  - barrier
    - synchronization mechanism that can be implemented using collective communication
  - reduction
    - combine collection of results from multiple UEs



# Design Spaces



## Recommended Reading

- Chapters 3, 4, 5, and 6 of  
Patterns for Parallel Programming  
by Mattson, Sanders, and Massingill
  - finding concurrency
  - algorithm structure
  - supporting structures
  - implementation mechanisms

