

**PMPP 2015/16**



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

# Performance Tuning (1)



# Course Schedule



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

---

12.10.2015	Introduction to PMPP
13.10.2015	Lecture CUDA Programming 1
19.10.2015	Lecture CUDA Programming 2
20.10.2015	Lecture CUDA Programming 3
26.10.2015	Lecture Parallel Basics, <a href="#">Exercise 1 assigned</a>
27.10.2015	Questions and Answers (Q&A), S3 19, Room 2.8
2.11.2015	<a href="#">Intro Final Proj.</a> , <a href="#">Ex. 1 due</a> , <a href="#">Ex. 2 assigned</a> , Lecture PRAM
3.11.2015	Lecture PRAM (2)
9.11.2015	<a href="#">Final Projects assigned</a> , L. Parallel Sort., <a href="#">Exercise 2 due</a>
10.11.2015	Questions and Answers (Q&A)
16.11.2015	Questions and Answers (Q&A)
17.11.2015	Questions and Answers (Q&A)
23.11.2015	<a href="#">1<sup>st</sup> Status Presentation Final Projects</a>
24.11.2015	<a href="#">1<sup>st</sup> Status Presentation Final Projects (continued)</a>
30.11.2015	Lecture Design Patterns
1.12.2015	Questions and Answers (Q&A)

---



# (Preliminary) Course Schedule

➔  
you are here

7.12.2015	Lecture Design Patterns (2), Performance Tuning
8.12.2015	Questions and Answers (Q&A)
14.12.2015	Performance Tuning (2)
15.12.2015	Questions and Answers (Q&A)
11.1.2016	2 <sup>nd</sup> Status Presentation Final Projects
12.1.2016	2 <sup>nd</sup> Status Presentation Final Projects (continued)
18.1.2016	Performance Tuning (3)
19.1.2016	
25.1.2016	
26.1.2016	
1.2.2016	
2.2.2016	
8.2.2016	Final Presentation Final Projects
9.2.2016	Final Presentation Final Projects (continued)



- **Instruction Level Parallelism**
- **Identifying Performance Limiters**
  - Instruction Limited Kernels
  - Latency Limited Kernels
  - Bandwidth Limited Kernels
  - Local Memory Optimizations
- **Multi GPU Multi Threading**
  - Brief review of the scenarios
  - Single CPU process, multiple GPUs
    - GPU selection, UVA, P2P
  - Multiple processes
    - Needs CPU-side message passing
  - Dual-IOH CPU systems and NUMA

Images and slides partially by [developer.nvidia.com](http://developer.nvidia.com)



# Instruction Level Parallelism (ILP)

- It is common to recommend
  - Running more threads per multiprocessor
  - Running more threads per thread block
- Motivation
  - This is the only way to hide latencies
- But...
- Two common fallacies
  - Multithreading is the only way to hide latency on a GPU
  - Shared memory is as fast as register



- Latency: Time required to perform an operation
  - $\approx 20$  cycles for arithmetic; 400+ cycles for memory
  - Can't start a *dependent* operation for this time
  - Can hide by overlapping with other code

**x** = a + b; // takes  $\approx 20$  cycles to execute

y = a + c; // independent, can start any time  
(stall)

z = **x** + d; // dependent, must wait for completion

# Arithmetic Throughput

- Latency is often confused with throughput
  - E.g. “arithmetic is 100x faster than memory – costs 4 cycles per warp (G80), whence memory operation costs 400 cycles”
    - One is rate, another is time
- Throughput: how many operations complete per cycle
  - Arithmetic: 1.3 Tflop/s = 480 ops/cycle (op = multiply-add)
  - Memory: 177 GB/s  $\approx$  32ops/cycle (op = 32-bit load)
- Little’s Law
  - Required Parallelism: Latency x Throughput



# Arithmetic Parallelism in Numbers

GPU model	Latency (cycles)	Throughput (cores/SM)	Parallelism (ops/SM)
G80-GT200	≈24	8	≈192
GF100	≈18	32	≈576
GF104	≈18	48	≈864
GF110	≈18	32	≈576
GK110	≈12	192	≈2304
GM200	≈8	128	≈1024


- Latency varies between different ops
- Can't get 100% throughput with less parallelism
  - Not enough operations in flight = idle cycles





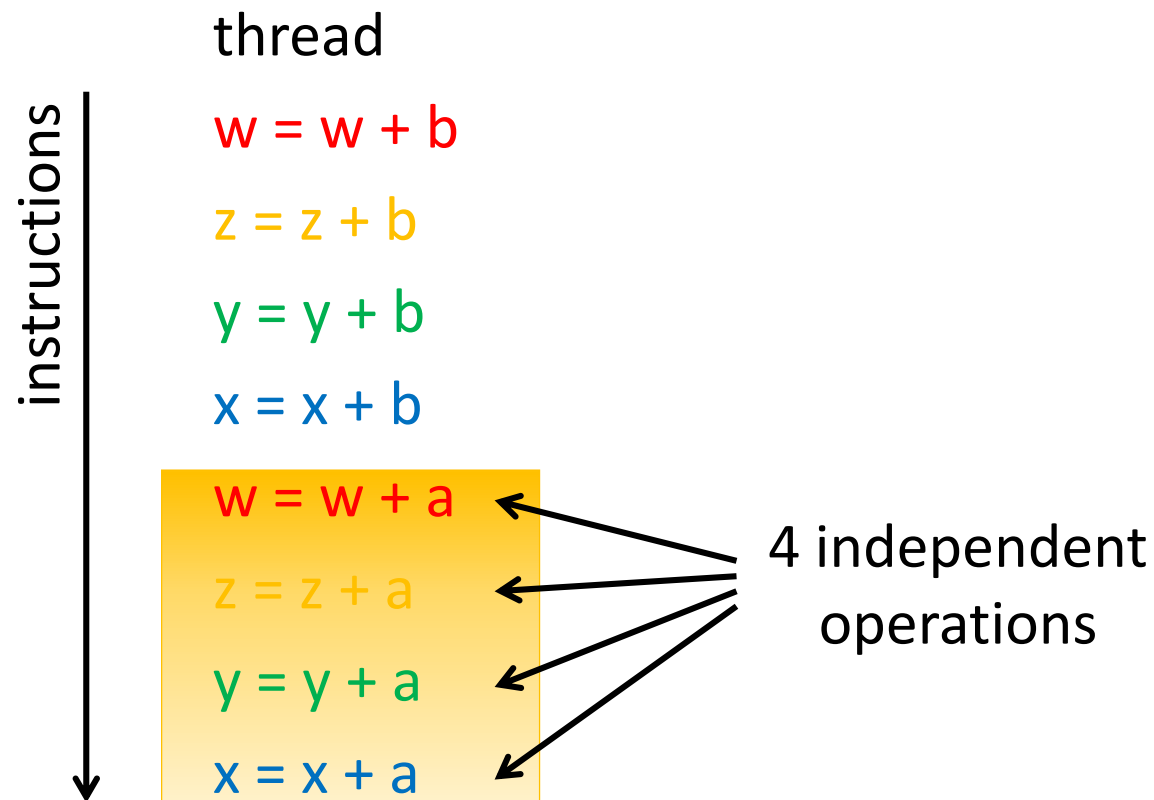
# Thread Level Parallelism (TLP)

- It is usually recommended to use threads to supply the needed parallelism, e.g. 192 threads per SM on G80:



# Instruction Level Parallelism (ILP)

- But you can also use parallelism among instructions in a single thread:



# You can use both ILP and TLP on GPU

- This applies to all CUDA-capable GPUs, e.g. G80:
  - Get  $\approx 100\%$  peak with 25% occupancy if no ILP
  - Or with 8% occupancy, if 3 operations from each thread can be concurrently processed
- On GF104 you *must* use ILP to get  $>66\%$  peak
  - 48 cores/SM, one instruction is broadcasted across 16 cores (half-warp)
  - So must issue 3 instructions per cycle
  - But have only 2 half-warp schedulers
  - Instead, it can issue 2 instructions per half-warp in the same cycle



# Let's check it experimentally

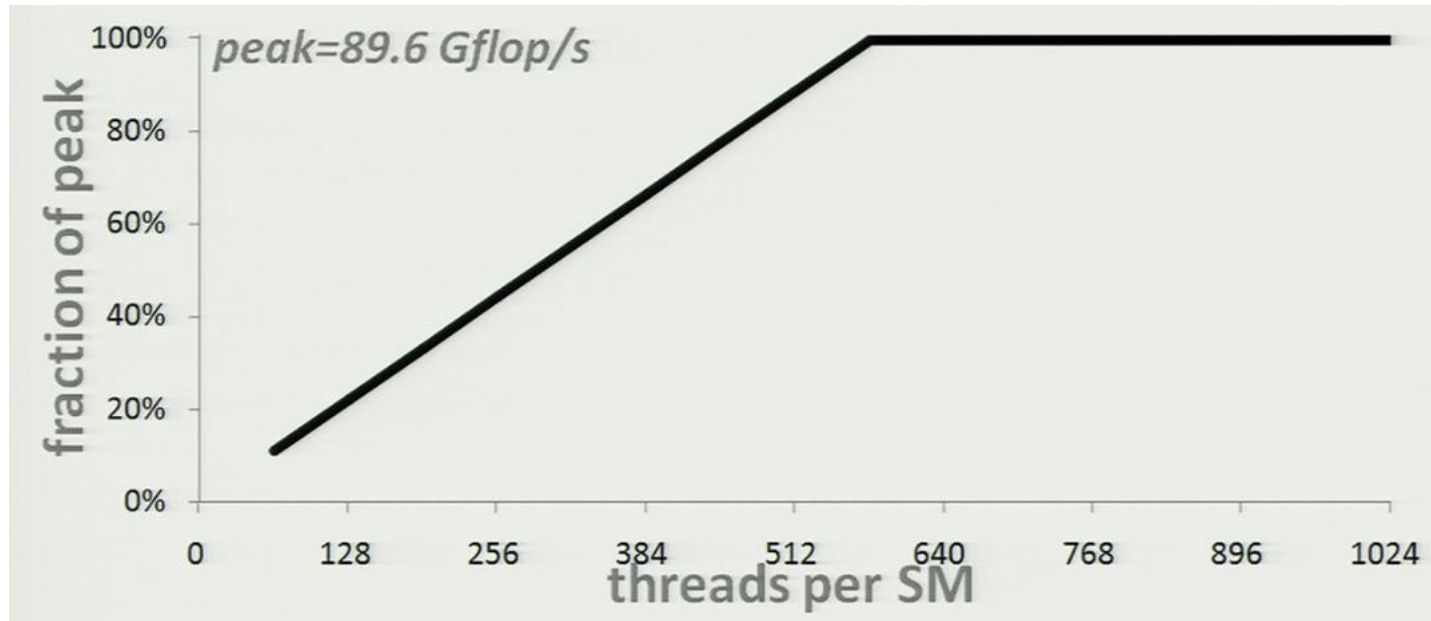
- Do many arithmetic instructions with no ILP:

```
#pragma unroll UNROLL
for (int i = 0; i < N_ITERATIONS; i++)
{
    a = a * b + c;
}
```

- Choose large **N\_ITERATIONS** and suitable **UNROLL**
- Ensure **a**, **b** and **c** are in registers and **a** is use later
- Run 1 block (use 1 SM), vary block size
  - See what fraction of peak (1.3TFLOPS/15) we get (gf100)



# Experimental results (GTX480)



- No ILP, need 576 threads to get 100% utilization

# Introduce Instruction Level Parallelism

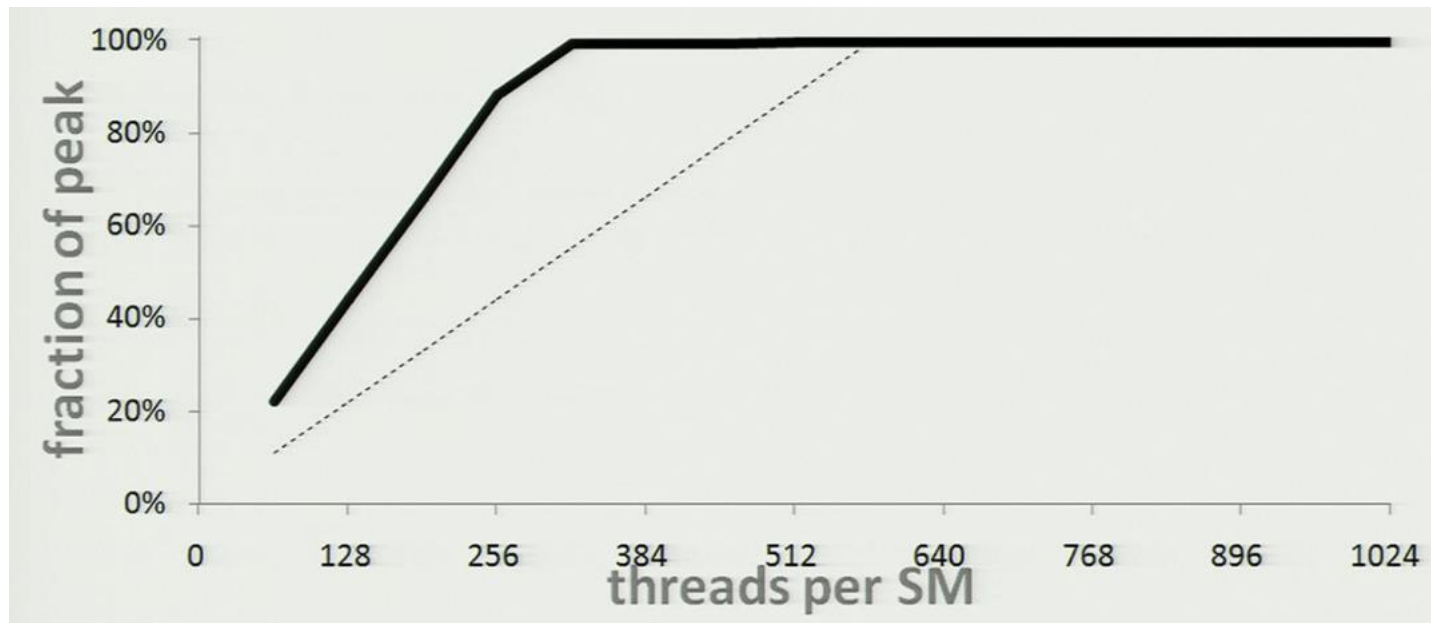
- Try ILP=2: two independent instructions per thread

```
#pragma unroll UNROLL
for (int i = 0; i < N_ITERATIONS; i++)
{
    a = a * b + c;
    d = d * b + c;
}
```

- If multithreading is the only way to hide latency on GPU, we've got to get the same performance



# GPUs can hide latency using ILP



- ILP=2: need 320 threads to get 100% utilization

# Add more Instruction Level Parallelism

- Try ILP=3: three independent instructions per thread

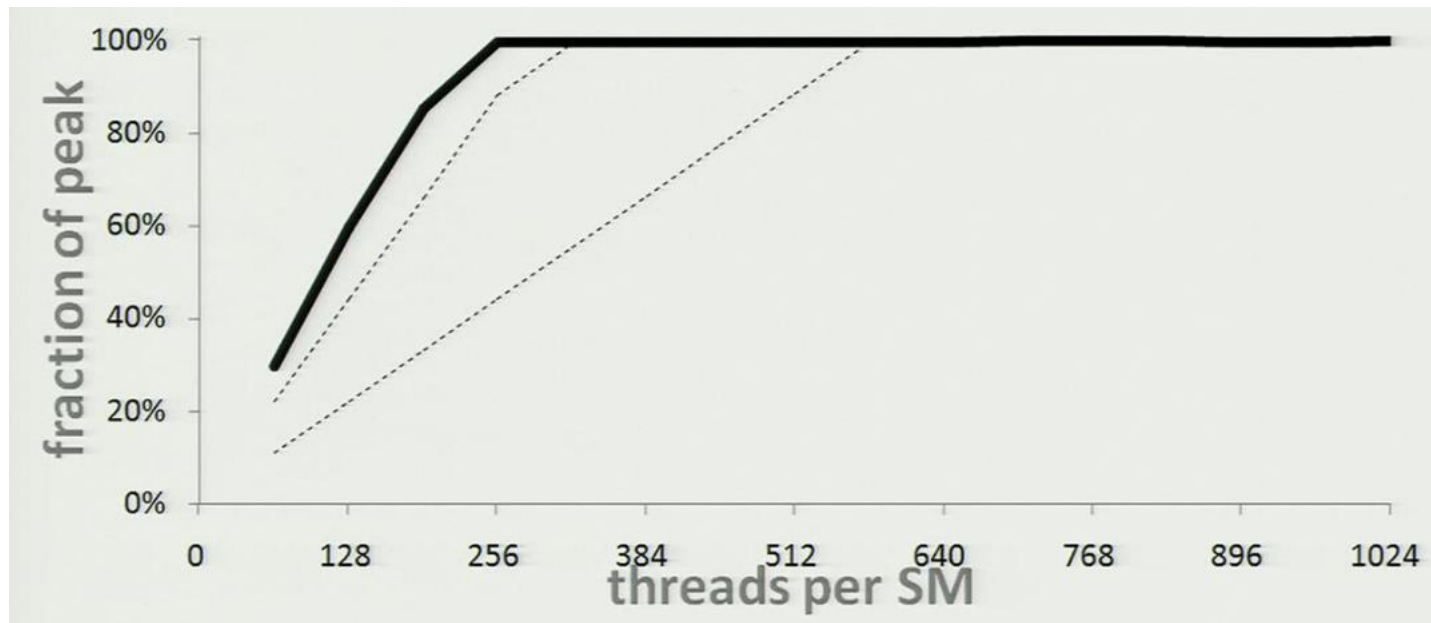
```
#pragma unroll UNROLL
for (int i = 0; i < N_ITERATIONS; i++)
{
    a = a * b + c;
    d = d * b + c;
    e = e * b + c;
}
```

- How far can we push it?



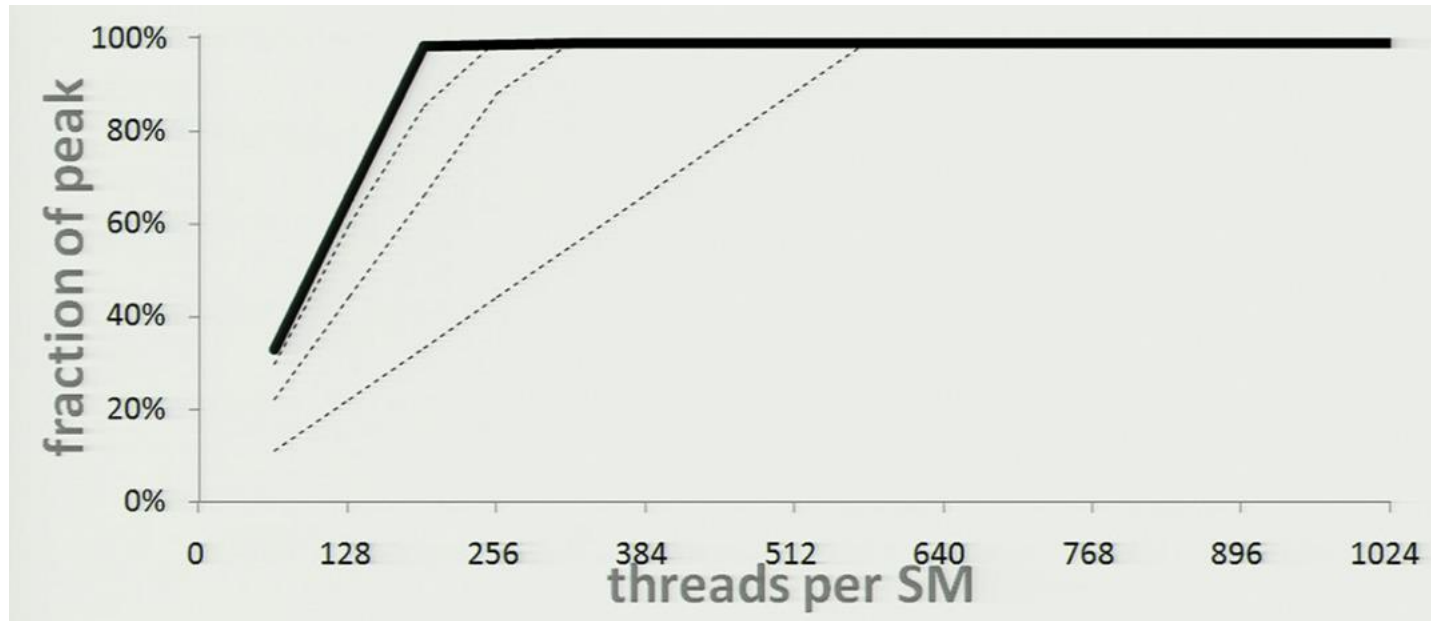


# Have more ILP – need fewer threads



- ILP=3: need 256 threads to get 100% utilization

# Unfortunately doesn't scale past ILP=4



- ILP=4: need 192 threads to get 100% utilization

# Hiding Memory Latency

- Apply same formula but for memory operations
  - Needed parallelism = Latency x Throughput

	Latency	Throughput	Parallelism
Arithmetic	≈18	32/SM/cycle	576 ops/SM
Memory	< 800 (?)	< 177 GB/s	< 100 KB

- So, hide memory latency = keep 100 KB in flight
  - Less if kernel is compute bound (needs fewer GB/s)



# How many threads is 100 KB?

- Again, there are multiple ways to hide latency
  - Use multiple threads to get 100 KB in flight
  - Use instruction parallelism (more fetches per thread)
  - Use bit-level parallelism (use 64/128-bit fetches)
- Do more work per thread – need fewer threads
  - Fetch 4B/thread – need 25,000 threads
  - Fetch 100B/thread – need 1,000 threads



- Copy one float per thread

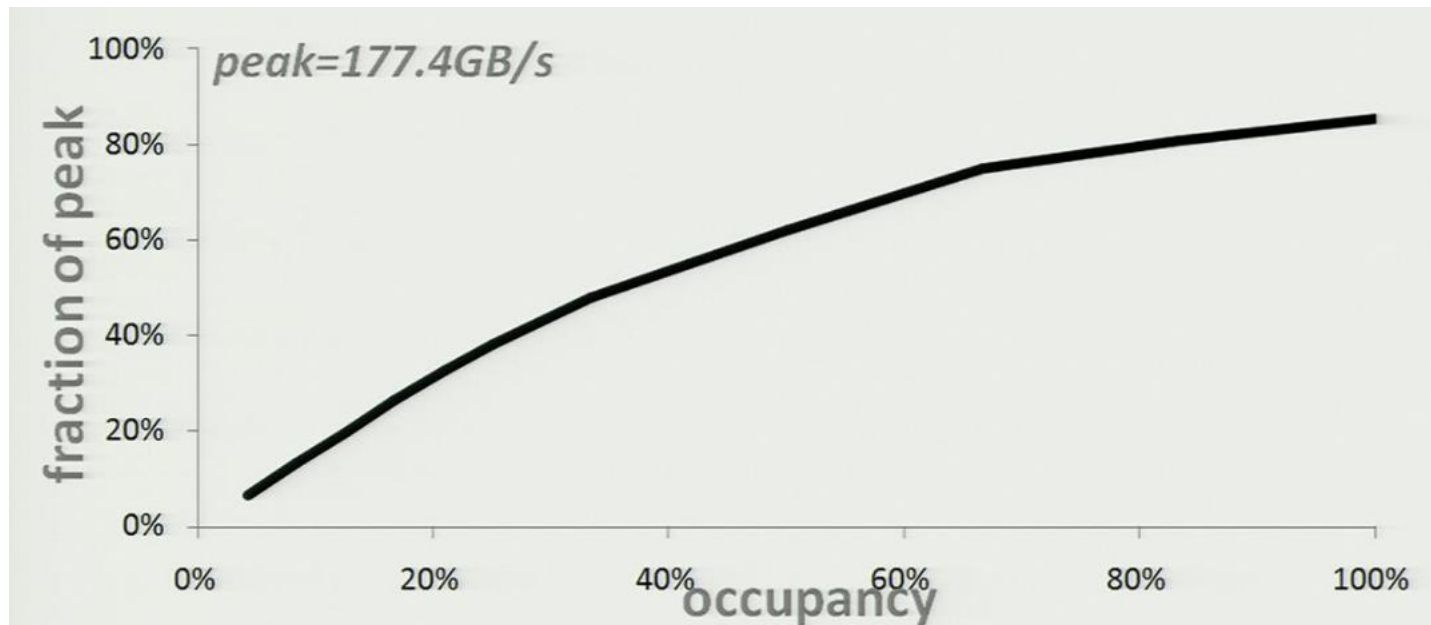
```
__global__ void memcpy(float *dst, float *src)
{
    int block = blockIdx.x + blockIdx.y * gridDim.x;
    int index = threadIdx.x + block * blockDim.x;

    float a0 = src[index];
    dst[index] = a0;
}
```

- Run many blocks, allocate shared memory dynamically to control occupancy



# Copying one float per thread (GTX480)



- Must maximize occupancy to hide latency?

# How far can we go?




- Copy 8 float4 per thread

```
__global__ void memcpy(float *dst, float *src)
{
    int block = blockIdx.x + blockIdx.y * gridDim.x;
    int index = threadIdx.x + block * blockDim.x;
    index = index * 8;

    float4 a[8];
    for (int i = 0; i < 8; i++)
        a[i] = src[index + i];
    for (int i = 0; i < 8; i++)
        dst[index + i] = a[i];
}
```




# Copying 8 float4 per thread (GTX480)



- 87% of peak bandwidth at only 8% occupancy!




# Copying 14 float4 per thread (GTX480)



- 84% of peak bandwidth at only 4% occupancy!

# Shared Memory Pitfall



# Running fast may require less Occupancy

- **Must use registers to run close to the peak**
- The larger the bandwidth gap, the more data must come from registers
- This may require many registers = **low occupancy**
- This often can be accomplished by computing multiple outputs per thread

